

# The Deadly Race

When Graceful Recovery Makes Coercion-Resistant Custody  
Worse Than Useless

*A formal threat model for physical-coercion (“wrench”) attacks,  
and a construction that clears it*

---

Yuri da Silva Villas Boas · Working paper, v2, September 2026

## Abstract

This paper contributes a *design bar* for coercion-resistant custody and a *construction* that meets it, motivated by a sharp negative result: custody designs that lack such a bar can be *worse than useless* — worse than merely a false sense of security, which would be bad enough — actively endangering the holder they claim to protect. Physical-coercion (“wrench”) attacks defeat custody schemes that key management and cryptography treat as out of scope. We give a *formal model* of coercion against a secret-holder and prove that a large class of “graceful” recovery designs — those that leave a device seizure holding a *feasible-but-unfinished* path to the secret — are strictly worse than useless: they manufacture a material incentive to *assassinate* the victim. We name this the **Deadly Race**, derive it from a single impossibility (one cannot prove one has forgotten), and distill a design bar — the **No-Gray-Area Criterion** — that any coercion-resistant scheme must clear. We then present a construction — the **Time-Gated Perceptual Oracle** (TGPO) — that clears the bar for the Bitcoin self-custody setting while keeping custody *individual*, and we locate a shipping alternative — timelock-rescue vaults — precisely: it clears the same bar by the *other* available route, at the cost of a trusted third party. We also classify the most widely shipped consumer countermeasure — the self-destruct (“wallet-erase”) PIN — and find it clears nothing: the vendor’s own documentation presupposes a surviving backup, so the erasure removes the device path while leaving intact the path that runs through the holder. The framework fixes a bar; it does not claim TGPO is the unique construction that clears it.

**Keywords:** Bitcoin; self-custody; wrench (physical-coercion) attacks; coercion resistance; Kerckhoffs’s principle; Deadly Race; No-Gray-Area Criterion; Tacit Knowledge-Based Authentication; Time-Gated Perceptual Oracle; memorization; hostage token; duress code; Obscurity Paradox; Denial Spiral; praxeology; shared custody; geographic custody.

## 1 Introduction

---

Coercion is the unmodeled attack. Standard threat models — Dolev–Yao on the wire [1], IND-CCA for encryption, forward secrecy for sessions — all assume the adversary cannot compel the *holder* of a secret. For bearer assets that assumption is empirically false: the physical attack registry maintained by Lopp [2], and the systematic study of the phenomenon in the financial-technologies literature [3], document a growing series of incidents in which a Bitcoin holder is physically threatened with a “\$5 wrench” until they surrender their keys.

Cryptographic hardness does not help here. A wrench attacker does not break the cipher; they compel the person who can operate it. Worse, the standard defense offered to holders — *obscurity* (“never tell anyone you own Bitcoin”) — is not a security mechanism at all in the sense of Kerckhoffs [4]: it is a bet that the attacker lacks knowledge the defender has, and it degrades precisely as adoption spreads that knowledge. A 2026 case already shows the degradation in miniature: a streamer who denied holding any Bitcoin was tortured for hours regardless, his denial simply disbelieved [5] (we return to it in §5.6). Nor is the leak channel usually the holder’s

own mouth: threat-intelligence reporting on the 2026 wave attributes victim selection to *leaked databases, tax records and exchange data*, alongside insider sales of customer lists and dark-web recruitment for employee access [6] — records held by third parties, which the holder can neither audit nor revoke once created.<sup>1</sup> This “keep a low profile, never disclose” counsel is nonetheless the community’s entrenched consensus, actively promoted by prominent voices (cataloged in the accompanying economic analysis [8]); and it carries a measurable price — that analysis finds the severity-adjusted actuarial cost of individual Bitcoin self-custody already *exceeds* the cost priced into insurance for comparable portable valuables such as jewellery.

This paper models coercion directly. Our contributions are:

1. a **formal coercion model** (§2) with an explicit adversary, a Kerckhoffs premise, and an engineering “feasibility line” that turns asymptotic security into concrete parameter bounds;
2. the **Deadly-Race impossibility chain** (§3): a classical impossibility (you cannot prove you forgot) that escalates, step by step, into an incentive to murder the victim;
3. the **No-Gray-Area Criterion** (§3), the design bar that a coercion-resistant scheme must clear, together with a *reachability* dichotomy that classifies every clearance route;
4. a **construction** (§4), the TGPO orbit, that meets the bar with individual custody, including a proof that the master secret meets the floor from every seizable state bar a bounded residual window; and
5. a **taxonomy** (§5) that classifies existing designs by *which* clearance route they take — illustrated on shipping time-lock-rescue vaults.

We call the central results the **Deadly-Race Lemma** and the **No-Gray-Area Criterion**, after their first statement in the protocol’s public design record [9].

## 2 Model

---

### Notation

The symbols below recur throughout; each is defined in full where it is introduced, and all are gathered here for reference.

---

<sup>1</sup>The counsel’s strongest form answers that such records need never exist: acquire only without identity verification, and keep no reportable trail. We grant that this reduces exposure and that it is legitimate personal hygiene. Three things nonetheless keep it from functioning as a security mechanism. *First, exposure ratchets.* A record cannot be uncreated and a leak cannot be un-leaked, so the strategy pays only *ex ante*, and only for the holder whose data has never once escaped; for everyone else there is no longer anything protectable to protect. Worse, the two curves run against the holder: the cost of maintaining discretion compounds with every year held and every counterparty met, while its marginal return falls as the stock of records already leaked grows. *Second, for the compliant holder the strategy is not lawfully available.* Where holdings or gains are legally reportable, leaving no fiscal record is not discretion but non-declaration — tax evasion, a criminal offense — so the counsel reduces to an instruction the law-abiding holder cannot follow. Nor is the resulting record safe by virtue of sitting with the state: in France, the jurisdiction carrying the largest share of recorded incidents, an employee of the tax administration was placed under judicial investigation in 2025 for allegedly consulting the authority’s internal database and passing the results to organized crime, with investigators reporting queries aimed specifically at cryptocurrency investors [7]. The compliant holder is thus asked to choose between an offense and a database that has been mined for targets. *Third, and decisive here,* discretion is not a *mechanism* at all: under Assumption 1 the adversary is granted the holder’s situation, so any defense contingent on their ignorance is out of scope by construction.

---

**Principals, assets, economics**


---

|                            |                                                                                                                                                                                   |
|----------------------------|-----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| $\mathcal{V}, \mathcal{A}$ | holder (victim) and coercer (attacker)                                                                                                                                            |
| $\mathcal{K}$              | <i>spend authority</i> : the set of states from which the stash can be spent; the gating condition may be a single key, a $k$ -of- $n$ multisig quorum, or a richer script (§3.1) |
| $W$                        | value of the asset to $\mathcal{V}$                                                                                                                                               |
| $c$                        | expected disutility to $\mathcal{A}$ of killing $\mathcal{V}$                                                                                                                     |
| $\pi$                      | $\mathcal{V}$ 's <i>pivotality</i> — win-probability $\mathcal{A}$ gains by removing $\mathcal{V}$                                                                                |

---

**Adversary, game, safety**


---

|       |                                                                                                                            |
|-------|----------------------------------------------------------------------------------------------------------------------------|
| WDY   | the Wrench–Dolev–Yao adversary, fixed by clauses D0–D8 (Assumption 2)                                                      |
| $p_A$ | probability $\mathcal{A}$ reaches $\mathcal{K}$ after the encounter; <i>gray</i> iff $0 < p_A < 1$                         |
| $t$   | coercion-time budget (how long $\mathcal{A}$ can sustain coercion)                                                         |
| $G$   | protocol time-gate (enforced latency to complete the derivation)                                                           |
| Corr  | residual-window event: the points $\mathcal{V}$ has surrendered so far are correct (§4.7)                                  |
| Esc   | residual-window event: the $\approx 1$ –2 min shaved by brute-forcing the last points decides escape from the scene (§4.7) |

---

**Construction (TGPO orbit §4)**


---

|                         |                                                                                                                   |
|-------------------------|-------------------------------------------------------------------------------------------------------------------|
| $\sigma$                | public, precomputation-resistant timechain salt (Namtso)                                                          |
| $H, H^*$                | unspecified cheap hash; unspecified memory- and depth-hard (Argon2-class) hash                                    |
| $o_i$                   | orbit state at stage $i \in \{0, \dots, N\}$ ; $o_i = H^*(u_{i-1})$ (with $o_0 = H(\sigma)$ ); $N+1$ stages total |
| $N$                     | number of <i>deep</i> stages — those requiring at least one memory- and depth-hard hash ( $H^*$ ) to derive       |
| $Sh_i$                  | Shamir polynomial at stage $i$ ; $r_i$ of its points reconstruct it                                               |
| $u_i$                   | $H(o_i \  Sh_i)$ : commitment of the wiped values $o_i$ and $Sh_i$                                                |
| $K_i$                   | $H(o_i \  Sh_i)$ : master secret correspondent to stage $i$ .                                                     |
| $p_{i,j}, \theta_{i,j}$ | point $j$ of stage $i$ ; parameters for fractal $j$ of stage $i$                                                  |
| $r_i$                   | Shamir threshold at stage $i$ ( $r_i \geq 2$ )                                                                    |
| $64^*$                  | $2^{64}$ memory- and depth-hard evaluations (infeasible)                                                          |
| floor                   | <i>prohibitive</i> cost: $\geq 64^*$ hard <i>or</i> $\geq 2^{96}$ cheap (Stipulation 1)                           |

---

## 2.1 Principals and capabilities

**Definition 1** (Principals). The **holder**  $\mathcal{V}$  possesses (in memory and on devices) whatever is needed to reach the spend authority  $\mathcal{K}$  gating an asset — the set of states from which it can be spent (§3.1); a single master secret is only the simplest case. The **coercer**  $\mathcal{A}$  seeks  $\mathcal{K}$ . Optional **heirs** or **delegates** may hold partial capability on  $\mathcal{V}$ 's behalf.

**Assumption 1** (Kerckhoffs, premise 0).  $\mathcal{A}$  knows the protocol in full: all algorithms, all source,  $\mathcal{V}$ 's setup topology, and every in-scope parameter. Only the protocol *secrets* are hidden. Every result below is proved *against* such an  $\mathcal{A}$ ; security by obscurity is assumed to contribute nothing.

**Assumption 2** (The Wrench–Dolev–Yao (WDY) adversary). The **Wrench–Dolev–Yao (WDY)** adversary inherits the Dolev–Yao [1] intruder *on the wire* — it sees all transmitted plaintext and ciphertext — but departs from it in two ways that matter here. It is *computational*, not symbolic: Dolev–Yao's perfect-cryptography idealization is replaced by the feasibility line (Stipulation 1), so “cannot break the cipher” becomes “cannot exceed its work budget”. And it is extended *past the wire* with the physical capabilities Dolev–Yao omits — device seizure and a bounded coercion oracle over the holder: the *wrench* the name records. WDY is fixed by:

**D0. Kerckhoffs.**  $\mathcal{A}$  knows the protocol,  $\mathcal{V}$ 's setup topology, and every in-scope parameter (Assumption 1).

- D1. Bounded compute.**  $\mathcal{A}$ 's computation is bounded by the feasibility line (Stipulation 1 below):  $2^{32}$  cheap evaluations feasible, plain  $2^{64}$  cheap at the border (the softest number), and  $2^{64}$  memory-hard ( $64^*$ ) and  $2^{96}$  cheap both infeasible. This is what replaces “unbounded”: the wall-clock custody window bounds  $\mathcal{A}$  crudely, the feasibility line *sharply*.
- D2. Frequency ceiling.** There is a ceiling on how much faster  $\mathcal{A}$ 's hardware evaluates a computation than  $\mathcal{V}$ 's. So an entropy-preserving hash with enough *non-parallelizable* steps ( $H^*$ ) lets  $\mathcal{V}$ 's setup impose an *arbitrarily long* latency on  $\mathcal{A}$  per evaluation — at the price of a proportionally longer latency on  $\mathcal{V}$ . This is what makes the time-gate meaningful:  $\mathcal{A}$  can neither parallelize nor out-clock its way past a sequential memory-hard chain.
- D3. Complete subjugation.** This is an information-security model, not a study of physical violence versus physical defense; we do not model  $\mathcal{A}$ 's coercive power against  $\mathcal{V}$ 's resistance. For simplicity, and to make the argument *a fortiori*, we assume that from the onset of the attack until  $\mathcal{A}$  releases  $\mathcal{V}$ ,  $\mathcal{A}$  meets *no resistance whatsoever* and enjoys  $\mathcal{V}$ 's *full cooperation*. The sole qualification — *feigned* cooperation in the residual window — is, as we argue there (§4.7), of residual importance.
- D4. Co-located device seizure.** All of  $\mathcal{V}$ 's devices *at the attack site* are immediately seized — vault blobs, time-lock blobs, commitment scripts, review schedules. A single WDY encounter has one *site*; devices and principals at *other* sites lie beyond its reach within the bounded window — the opening a *geographic* defense exploits (§3.6).
- D5. Explicit-secret seizure.** All of  $\mathcal{V}$ 's *explicit* (non-tacit) authenticating secrets are immediately seized.
- D6. No tacit seizure.**  $\mathcal{V}$ 's *tacit* entropy-carrying secrets — the points — are captured by seizure *if and only if* they are *available* in transmissible form at the moment of attack (materialized, and not wiped before  $\mathcal{V}$  is subdued) — that is, exactly when the tacitness fails to be effective. When it *is* effective they are not: by D3 this is a genuine *incapacity to transmit*, not any unwillingness — even a fully cooperating  $\mathcal{V}$  cannot hand them over, because they are recognizable but non-transmissible. This and D7 are the operational restatement, in the adversary model, of the tacit-knowledge premise (Assumption 3).
- D7. No perceptual-parameter seizure.** As with the tacit entropy-carrying secrets (D6),  $\mathcal{V}$ 's perceptual-oracle parameters — the deep fractals — are seized *iff* they are available at the moment of attack; when the tacitness is effective they are not. As in D6, and by D3, this is incapacity under full cooperation, not withholding: the deep fractal parameters are regenerated only by walking the memory-hard chain, so a willing  $\mathcal{V}$  has nothing to dictate (Assumption 3).
- D8. No remote-delegate coercion.** A remote delegate — a party holding rescue capability on  $\mathcal{V}$ 's behalf beyond the attack site (D4) — is past  $\mathcal{A}$ 's *remote* coercive reach: from the encounter with  $\mathcal{V}$ ,  $\mathcal{A}$  cannot compel it to surrender or act on its state. Without this a delegated defense is vacuous. Real delegates span a spectrum of coercion-compliance; collapsing it to a clean “cannot be reached” is a deliberate simplification — the kind that makes the *a fortiori* argument possible, at the cost of reality's shades.

The *only* channel to the tacit material withheld by D6–D7 is a **bounded coercion oracle**:  $\mathcal{A}$  extracts it through the cooperating  $\mathcal{V}$  (D3) over a bounded duration  $t$  — the adversary's **coercion-time budget**, the maximum wall-clock time for which it can sustain coercion, calibrated against the incident durations in the registry [2]. Cooperation is genuine throughout by D3; the one exception, *feigned* cooperation, is confined to the residual window (§4.7).

*Remark 1* (The two non-seizure clauses coincide for design). For protocol-design purposes the tacitness of the entropy-carrying secrets (D6) and that of the perceptual-oracle parameters (D7) are logically equivalent. If the parameters were transmissible, a wrench attack would seize them — and, available within that *same* encounter, they would then serve to seize the entropy-carrying points themselves; so non-tacit parameters force non-tacit points. Conversely, if the points were somehow transmissible without the parameters ever being available, then the tacitness of the

parameters is irrelevant, the points having already fallen. Either way the binary question “does a single seizure reach  $\mathcal{K}$ ?” turns on the two clauses together, and the model may treat them as one.

*Remark 2* (Where the distinction still matters). The distinction can nonetheless matter for a *quantitative* refinement — a statistical distribution of the wall-clock needed to coerce either (up to sufficient number of most-significant bits) given  $\mathcal{A}$ ’s coercion time budget  $t$  (Definition 4). Still, a secret that is *theoretically* transmissible, but only with pen and paper, an adversarial fork of the application, and long hours to days of focused mental effort under perfect  $\mathcal{A}$ – $\mathcal{V}$  cooperation, is essentially as good as ideal. Until empirical studies of TKBA (Assumption 3) are done, the actual relevance of that nuance is purely speculative.

## 2.2 The feasibility line

Asymptotic “hard” and “easy” are useless against a fixed, resourced attacker; we replace them with two engineering stipulations that fix what *feasible* means and turn the lemmas into concrete bounds.

**Stipulation 1** (Feasibility line). Let a *cheap* evaluation be one call to a non-memory-hard hash  $H$ , and a *hard* evaluation one call to an hours-scale Argon2-class function  $H^*$  whose cost is a large number of *non-parallelizable* steps — *memory- and depth-hard*,<sup>2</sup> so  $\mathcal{A}$ ’s bounded speed advantage (D2) cannot amortize it [11].

**S1 (feasible).**  $2^{32}$  cheap evaluations are within  $\mathcal{A}$ ’s budget.

**S2 (infeasible).**  $2^{64}$  *hard* evaluations — written  $64^*$  — are beyond it.

**S3 (infeasible).**  $2^{96}$  cheap evaluations are beyond it.

The star denotes this memory- and depth-hardness and attaches to *any* bit-count —  $32^*, 64^*, 96^*, \dots$  — each meaning that many  $H^*$  evaluations, not just  $64^*$ . Neither S1 nor S2 settles the star at 32 bits, so we state it:  $32^*$  **is feasible** — memory-hardness alone does not lift a single 32-bit stage over the floor, which is why no orbit link commits fewer than 64 bits (§4).

**The floor.** We call a work cost *prohibitive* — equivalently, say it *meets the floor* — when it satisfies S2 *or* S3: at least  $64^*$  memory-hard evaluations, *or* at least  $2^{96}$  cheap ones. The two bounds are independent, and we deliberately fix *no* ordering between them. For realistic memory- and depth-hardness the per-evaluation \* premium is well under  $2^{32}$ , so  $64^* < 2^{96}$  and the cheap bound is the binding one; but \* is a user-tunable lever, and an aggressive enough setting pushes that premium past  $2^{32} \approx 4.29 \times 10^9$  — above 32 extra bits — making  $64^* \geq 2^{96}$ . Every “floor” claim below names *this* threshold, not a fixed number, and holds whichever bound binds.

We keep the star rather than fold it into an equivalent bit-count because memory- and depth-hardness has *three* distinct consequences, only the first of which a bit-count could absorb: it (1) inflates the brute-force cost — this alone *could* be traded for a stipulated increase in bits; (2) imposes a wall-clock *derivation time* on each evaluation, which the bounded coercion window  $t$  (D3) cannot outrun; and (3) rules out the quantum (Grover) speed-up (§6). A pure bit-count captures only (1); the star carries all three.

Two currencies span the gap between them. *Entropy* buys  $32 \rightarrow 64$ : a second independent 32-bit stage. But plain  $2^{64}$  *cheap* work still sits at the border of a resourced attacker’s reach (hours), so the  $64 \rightarrow 64^*$  step must be bought by *memory-hardness* — which is why the gating function has to be Argon2-class. We refer to plain cheap  $2^{64}$  throughout as the *softest number*: it must be named explicitly wherever it is invoked, never silently treated as infeasible.

<sup>2</sup>*Depth-hard* is meant more sharply than *sequential* or *non-parallelizable* alone: the cost is a non-parallelizable computation with a *deliberately large number of steps* — a long sequential *depth* — the parallelism-resistance axis of *depth-robust* memory-hard functions [10] — that  $\mathcal{A}$  can neither parallelize away nor out-clock past (D2). We also weighed *memory- and sequentiality-hard* (accurate but clumsy) and *memory- and iteration-count-hard* (which names the step *count* rather than its non-parallelizable structure); *depth* is the single noun that carries both the seriality and the magnitude, and it parallels *memory*.

### 2.3 The tacit-knowledge premise

The construction of §4 rests the secret on *tacit* knowledge. We state the premise abstractly here; §4 instantiates it.

**Assumption 3** (Tacit knowledge, TKBA). The secret is a set of fractal locations that are *recognizable but not transmissible*.

- **TKBA<sub>1</sub>**. A point is non-transmissible without the fractal parameters that render it;  $\mathcal{V}$  cannot dictate it in the abstract.
- **TKBA<sub>2</sub>**. The *deep* fractal parameters are not a hand-overable artifact; producing them requires walking a memory-hard chain, before which later locations do not exist.

*Remark 3* (TKBA restates the tacit-seizure clauses). Assumption 3 is not an independent hypothesis: it is clauses D6–D7 lifted out of the adversary model and stated as a premise on the secret. The clauses say  $\mathcal{A}$  cannot seize the tacit points or the perceptual parameters even under full cooperation (D3); TKBA says the same thing about the secret itself — recognizable but not transmissible — with no adversary mentioned. They are two faces of one assumption (exactly as D0 and Assumption 1 are two faces of Kerckhoffs), and by Remark 1 the two clauses are themselves one. We keep both forms because each is the natural one in its place — the premise where the construction is motivated (§4), the adversary clause where seizure is reasoned about.

## 3 The Deadly-Race impossibility chain

*“Nothing was your own except the few cubic centimetres inside your skull.”*  
— George Orwell, *Nineteen Eighty-Four*

This section states the paper’s core negative result as an *escalating chain*: the *Non-Negative-Proof Principle*, an *a fortiori* assumption on the attacker’s knowledge (justified by an information-theoretic impossibility, §3.2), its decision-theoretic consequence for a rational coercer (§3.3), and the design bar the two together force (§3.4–§3.5). We are deliberately explicit about which links are *impossibilities* and which are *incentive* claims: the Deadly-Race Lemma is **not** a cryptographic hardness statement, and reading it as one would be exactly the overclaim §8 warns against.

### 3.1 Setup and definitions

Fix the model of §2. Write  $\mathcal{K}$  for the *spend authority* protecting an asset worth  $W$  to the holder  $\mathcal{V}$ , and  $c$  for the expected cost to  $\mathcal{A}$  of killing  $\mathcal{V}$  — the legal, moral, and operational disutility of a homicide. Bitcoin makes the prize *rivalrous* and *bearer*: the *first* party to broadcast a valid spend takes  $W$ ; there is exactly one prize and no appeal.

$\mathcal{K}$  is a *set of states* — states in the same sense as the seizable state  $\Sigma$  of Definition 2 — namely those from which the stash can be spent; *reaching*  $\mathcal{K}$  abbreviates *reaching some element of*  $\mathcal{K}$ . Nothing in what follows requires  $\mathcal{K}$  to be the singleton “knows one private key”. A single-key address is the special case; a  $k$ -of- $n$  multisig makes  $\mathcal{K}$  the states holding any  $k$  of the  $n$  shares, and script or covenant conditions extend it further. The distinction is load-bearing exactly where the analysis turns on *which* pieces a party must hold (§3.6, §3.7); everywhere else the reader may safely picture a key.

**Definition 2** (Seizable state, feasible path, racers).

- **Seizable state**  $\Sigma$ . Everything  $\mathcal{A}$  obtains from the encounter — all seized device state, plus whatever  $\mathcal{A}$  extracts from  $\mathcal{V}$  during the bounded coercion window.
- **Feasible path**. A *recovery path* from a state is an algorithm that outputs an element of  $\mathcal{K}$  within the feasibility line (Stipulation 1). A state *feasibly reaches*  $\mathcal{K}$  if such a path exists; otherwise reaching  $\mathcal{K}$  from it is *infeasible*.




Figure 1: **The Deadly Race decision.** After a WDY seizure (Assumption 2), the outcome is fixed by  $p_A$ , the probability  $\mathcal{A}$  reaches  $\mathcal{K}$ . At the *clear* endpoints ( $p_A \in \{0, 1\}$ ) no race exists. In the *gray* interval  $0 < p_A < 1$  the holder  $\mathcal{V}$  is *pivotal* ( $\pi > 0$ , Theorem 1): removing  $\mathcal{V}$  raises  $\mathcal{A}$ 's win-probability, so once the asset clears  $W > c/\pi$  a rational  $\mathcal{A}$  has a positive incentive to *assassinate* (Lemma 1). Clearing the No-Gray-Area Criterion (Criterion 1) means removing the gray branch entirely.

- **Post-release racers.** If  $\mathcal{A}$  releases  $\mathcal{V}$ , two parties may pursue  $\mathcal{K}$ :  $\mathcal{A}$  from  $\Sigma$ , and  $\mathcal{V}$  from memory / tacit recall. Let  $p_A \in [0, 1]$  be the probability that  $\mathcal{A}$  broadcasts a valid spend *before*  $\mathcal{V}$  does, given the states each holds after release. Because the prize is rivalrous,  $\mathcal{V}$ 's countermove — spending first, to safety — is  $\mathcal{A}$ 's only loss condition once  $\mathcal{A}$  holds *any* feasible path.

**Definition 3** (Gray vs. clear outcomes). A coercion outcome is **gray** if afterward  $0 < p_A < 1$ :  $\mathcal{A}$  has a feasible path ( $p_A > 0$ ) yet  $\mathcal{V}$ , alive and free, can still win the race ( $p_A < 1$ ). The **clear** outcomes are the endpoints —  $p_A = 1$  ( $\mathcal{A}$  certainly wins) and  $p_A = 0$  ( $\mathcal{A}$  has no feasible path at all). Figure 1 traces the decision this induces.

Stated on the encounter rather than on the racers: gray is where the encounter leaves  $\mathcal{A}$  a state *sufficient* for  $\mathcal{K}$  but does not complete the spend — the derivation or time-lock outstanding at release exceeds what remains of the coercion budget  $t$  (D3). We therefore call such a path **feasible-but-unfinished**, not *feasible-but-slower*: “slower” names a comparison with  $\mathcal{V}$  that  $\mathcal{A}$  can never settle, since by NNPP (Assumption 4) an undisclosed shortcut may exist, so  $p_A$  does not reach 1 while  $\mathcal{V}$  is alive and free. Pivotality then follows from the encounter alone — absent an access-dominating delegate (Definition 10), whose presence leaves the outcome gray but *safe*, the pivotal racer being out of  $\mathcal{A}$ 's reach.

**Definition 4** (Priced and sharpened adversaries). The **priced** WDY adversary, written  $\text{WDY}[W, c]$ , is a WDY adversary (Assumption 2) additionally endowed with the economic parameters  $(W, c)$  and modeled as a *cold expected-utility maximizer*: it chooses among release, continued coercion, and assassination so as to maximize its assessed expected payoff over the deadly-race probabilities (the win probability  $p_A$ , and the correctness/escape events introduced in §4.7). Plain WDY carries *no* such payoff: it is the adversary of the general results (NNPP, the No-Gray-Area Criterion, the floor), which must hold whatever  $(W, c)$  turn out to be.  $\text{WDY}[W, c]$  is invoked only where a *quantitative* incentive is actually computed — the Deadly-Race Lemma's condition (§3.3) and the residual-window analysis (§4.7).

A third, optional parameter sharpens the feasibility line to a single threshold.  $\text{WDY}[W, c, b]$  replaces the two stipulations of Stipulation 1 by one *clear-cut* bound  $b$  — total work below  $b$  feasible, at or above  $b$  infeasible — in cost-homogeneous units, with memory-hardness folded into the per-evaluation cost. This idealization is legitimate here because the protocol quantizes entropy in 32-bit chunks, so feasibility is a step function on the lattice  $\{32, 64, 96, \dots\}$  and its hard counterparts  $\{32^*, 64^*, 96^*, \dots\}$ : a threshold at that granularity loses nothing, whereas pricing a single

brute attempt would be finer than any distinction the construction can express. Conservatively  $b = 64^*$  — work below  $64^*$  feasible,  $64^*$  itself infeasible. A fourth optional parameter  $t$  budgets the **coercion time** (D3’s window):  $\text{WDY}[W, c, b, t]$  fixes the maximum wall-clock time for which the adversary can sustain coercion. Like  $b$ ,  $t$  is a capability bound pinned at its real, registry-calibrated value  $t_{\text{reg}}$ , not sent to  $\infty$  — unbounded coercion would let  $\mathcal{A}$  walk any time-gate to completion, the same category error as  $b \rightarrow \infty$ . The rungs are thus  $\text{WDY}$  (capabilities only),  $\text{WDY}[W, c]$  (+ payoff),  $\text{WDY}[W, c, b]$  (+ sharp feasibility), and  $\text{WDY}[W, c, b, t]$  (+ coercion-time budget).

*Remark 4* (Plain  $\text{WDY}$  as a worst case). Plain  $\text{WDY}$  is exactly the worst case of the priced family over the *payoff* coordinates, taken at the model’s *fixed* capability:

$$\text{WDY} = \text{WDY}[W \rightarrow \infty, c \rightarrow 0, b = 64^*, t = t_{\text{reg}}].$$

The payoff coordinates are sent to their extreme — infinite stake, costless homicide: the maximally kill-prone attacker, which assassinates in *every* gray outcome (§3.3) — because the general results must hold whatever  $(W, c)$  are. The capability coordinates are **not** worst-cased:  $b$  stays pinned at the feasibility line  $64^*$  and  $t$  at the registry window  $t_{\text{reg}}$ . Sending  $b \rightarrow \infty$  would grant  $\mathcal{A}$  a  $64^*$  brute and collapse the floor (Theorem 3); sending  $t \rightarrow \infty$  would let it walk any time-gate to completion and collapse the tacit route (§3.6). Either is a category error, not a stronger guarantee. The asymmetry has a tell at the residual window: there  $p_A \rightarrow 1$  while the worst case wants  $W \rightarrow \infty$ , making  $(1 - p_A)W$  an  $\infty \cdot 0$  indeterminate form — which is precisely why that window is settled with *finite*  $(W, c)$  and the  $\Delta p_A$  increment (§4.7), not with the blunt limit.

### 3.2 The Non-Negative-Proof Principle

**Assumption 4** (Non-Negative-Proof Principle, NNPP — *a fortiori* on Kerckhoffs). Kerckhoffs (Assumption 1) grants  $\mathcal{A}$  the *protocol* — and that  $\mathcal{V}$  *may* hold brain-memory or an undisclosed backup short-cutting  $\mathcal{K}$  — but *not* how many such backups  $\mathcal{V}$  keeps, or where. NNPP covers exactly that gap: what  $\mathcal{A}$  *cannot* know is whether any undisclosed backup *actually* exists. The two do not conflict — one is the public mechanism, the other  $\mathcal{V}$ ’s private state — so NNPP *qualifies* Kerckhoffs, and we resolve the gap conservatively:  **$\mathcal{A}$  always assumes an undisclosed backup exists**. Notice that something as small as an SD card, as inconspicuous as steganographic notes in a book, or as interior as the memorized login and password of an email account suffices to store a short-cutting state.

Both halves point the *same* way — toward killing. Knowing the protocol (Kerckhoffs) and always presuming  $\mathcal{V}$  a live racer (NNPP) each *raise*  $\mathcal{A}$ ’s assessed deadly-race payoff  $\pi W$ , hence its proclivity to kill — precisely the outcome the method must survive. Each is therefore *a fortiori*: we endow  $\mathcal{A}$  with the *most kill-prone* epistemics and demand safety even then. In one line:

*if, even after maximizing the attacker’s proclivity to enact the deadly race, killing still buys it nothing, the method is (a fortiori) deadly-race-safe.*

To avoid a nightmarishly complex bluffing game between  $\mathcal{A}$  and  $\mathcal{V}$ , and to avoid encouraging real users to bet their stashes — and their lives — on it (see Remark 33), we also do not consider  $\mathcal{A}$  reading truthfulness off  $\mathcal{V}$ ’s terrorized demeanor or tone. The simplification this forces is that  $\mathcal{A}$  is *uniformly distrusting*: it always defaults to assuming that whatever NNPP posits *might* exist does in fact exist. The one alternative simplification — a *gullible*  $\mathcal{A}$  that always believes  $\mathcal{V}$  — is trivially sub-NNPP and sub-Kerckhoffs, hence useless as a threat model. Erring toward distrust is therefore the only valid choice, and it is what makes the model *a fortiori*: a protocol that delivers seizure resistance or deadly-race resistance against a uniformly distrusting attacker delivers them against a real one too.

In practice NNPP reduces to a plain observation: **exhaustive search is so ineffective that  $\mathcal{A}$  always assumes another backup exists, whatever the terrorized  $\mathcal{V}$  pleads**. It is

rational, not paranoid — non-possession is a *negative existential* (“nothing anywhere in  $\mathcal{V}$ ’s reach reproduces  $\mathcal{K}$ ”) that cannot be certified on either reading of “prove.”

*Remark 5* (Why even total search fails — the self-defeat). The search reading fails not merely because the covering is large but because the capability that would complete it is *off the domain of the problem*. To certify  $\neg\exists$  witness,  $\mathcal{A}$  must reach the last, irreducible site — the “few cubic centimetres” of the epigraph: whether  $\mathcal{V}$  has *forgotten*, or never committed,  $\mathcal{K}$  to memory. An  $\mathcal{A}$  that can read that space, and cover all physical hiding-space besides, is *supra-totalitarian* — it exceeds the literary supremum of surveillance power, which expressly conceded the skull — and it is a *degenerate fixed point* of the model: it violates the bounded-adversary premise so completely that the contested-property premise collapses with it, for under an uncontested overlord there is no custody to resist. Defending against it is a *doomsday bet*: the branch in which the adversary exists is the branch in which the stakes that define the game are gone, so a rational actor assigns it *zero decision-weight* — not because its probability is negligible, but because its utility is undefined there. NNPP is therefore safe as a *modeling boundary*: the only world that falsifies it is one in which the question does not arise.

**Corollary 1** (Cooperation cannot dispel the incentive). *Let  $b$  be the number of resumable copies  $\mathcal{V}$  holds, and under the coercion oracle ( $D3$ , Assumption 2) let  $\mathcal{V}$  disclose some of them. Each disclosure confirms a specific copy existed, raising a known lower bound  $\hat{b} \leq b$ ; none yields an upper bound  $\hat{b} \geq b$ , which would require the covering search of Remark 5. Disclosure is thus a monotone lower-bound estimator — “ $\mathcal{V}$  has revealed everything” is precisely the direction NNPP forbids. Hence no degree of cooperation drives  $\mathcal{A}$ ’s posterior on “ $\geq 1$  undisclosed copy” to zero: in any design that leaves a released  $\mathcal{V}$  some resumable path, the assassination incentive of the Deadly-Race Lemma (§3.3) cannot be discharged by surrender — only by the victim’s elimination.*

### 3.3 The Deadly-Race Lemma

**Lemma 1** (Deadly-Race Lemma). *Let a custody scheme admit a gray coercion outcome (Definition 3): on release,  $\mathcal{A}$  holds a feasible path to  $\mathcal{K}$  ( $p_A > 0$ ) while  $\mathcal{V}$  retains one too ( $p_A < 1$ ). Then the priced adversary  $\text{WDY}[W, c]$  (Definition 4) has a **strictly positive incentive to eliminate**  $\mathcal{V}$  (and any co-racing heirs) whenever  $(1 - p_A)W > c$ .*

*Proof.* With  $\mathcal{V}$  alive and free,  $\mathcal{A}$ ’s expected take from the race is  $p_A W$ . By NNPP (Assumption 4),  $\mathcal{A}$  cannot instead obtain or demand a verified assurance that  $\mathcal{V}$  will not race — no such proof exists — so the *only* instrument available to  $\mathcal{A}$  for pushing the win probability toward 1 is to remove the competing racer. Eliminating  $\mathcal{V}$  drives  $\mathcal{A}$ ’s win probability to  $\approx 1$  and the expected take to  $\approx W$ . The marginal gain is  $W - p_A W = (1 - p_A)W$ , and a rational  $\mathcal{A}$  acts on it exactly when it exceeds the cost  $c$ . In any gray outcome  $p_A < 1$ , so  $(1 - p_A)W > 0$ ; for assets large relative to  $c$  the inequality holds.  $\square$

Three remarks fix the result’s meaning and scope:

*Remark 6* (It is decision-theoretic, not cryptographic). Lemma 1 says nothing about breaking a cipher; it holds *whatever* the underlying cryptography, because it lives in the incentive layer threat models omit. This is why “the crypto is strong” is no defense — the danger is created by the *shape* of the recovery path, not its hardness.

*Remark 7* (Release is the trigger). The incentive is a property of the released gray state: it is  $\mathcal{V}$ -as-free-racer whom  $\mathcal{A}$  is paid to remove. A scheme that simply hands  $\mathcal{A}$  the funds outright kills no one; the gray fallback is *worse than* that clean loss. And by Corollary 1 no amount of surrender removes the incentive —  $\mathcal{A}$  can never confirm  $\mathcal{V}$  kept no resumable copy — so only the victim’s elimination closes it.

*Remark 8* (The danger peaks where the marketing promises safety). A “graceful” fallback — a slow-but-feasible recovery left on a seized device, a time-locked rescue, a delayed vault — holds  $p_A$

well below 1 while leaving  $W$  fully at stake, i.e. it *maximizes*  $(1 - p_A)W$ . Designs that degrade gracefully thus tend to sit squarely in the region of **largest** assassination incentive. Graceful degradation, a virtue everywhere else, is here a liability — hence the title.

### 3.4 The No-Gray-Area Criterion

**Criterion 1** (No-Gray-Area Criterion). *Under a fixed threat model, a coercion-safe scheme must admit **no gray outcome**: every encounter must resolve to a clear endpoint —  $p_A = 0$  (the attack plainly fails:  $\Sigma$  feasibly reaches nothing and  $\mathcal{V}$  holds nothing transmissible left to surrender) or  $p_A = 1$  (it plainly succeeds:  $\mathcal{A}$  obtains  $\mathcal{K}$  within the encounter, so no post-release race exists). The forbidden interval  $0 < p_A < 1$  is exactly where Lemma 1’s incentive lives.*

The Criterion is a **bar, not a construction**: it fixes what a design must achieve (collapse every outcome out of the gray interval), not how. The two clear endpoints are not equally desirable to  $\mathcal{V}$  —  $p_A = 1$  costs  $\mathcal{V}$  the funds — but both are *survivable*: neither leaves a released racer whom  $\mathcal{A}$  profits by killing. Coercion resistance, properly posed, is therefore a bar on **the holder’s safety** first, and only then on the funds.

#### The Criterion as a theorem

The Criterion need not stay a stipulation. Cast the encounter as a one-move game and “coercion-safe” as a predicate on it, and no-gray becomes exactly the safety condition.

**Definition 5** (Coercion game and coercion-safety). Fix a scheme  $\Pi$  and a priced adversary  $\text{WDY}[W, c]$  (Definition 4). The game **Coerce** $_{\Pi}$  is: (1)  $\mathcal{A}$  mounts an encounter — device seizure (D3–D7) and the bounded coercion oracle (D3) — reaching a post-encounter state with a set  $R$  of *racers*: the parties that can still reach  $\mathcal{K}$  within feasibility and act on it (spend, or rescue/cancel to defeat  $\mathcal{A}$ ’s spend) — the released  $\mathcal{V}$ , and any remote delegate. Write  $p_A(S)$  for  $\mathcal{A}$ ’s probability of prevailing against racer set  $S$  (monotone: fewer racers, weakly higher  $p_A$ ); the scalar  $p_A$  of Definition 2 is  $p_A(R)$ . (2)  $\mathcal{A}$  picks  $a \in \{\text{spare}, \text{kill}\}$ , where kill removes  $\mathcal{V}$  from  $R$  — by reachability  $\mathcal{A}$  can eliminate only racers it can *reach*, never a remote delegate. (3) the race resolves, with payoff

$$U_{\mathcal{A}}(a) = p_A(S_a) \cdot W - c \cdot \mathbf{1}[a = \text{kill}], \quad S_{\text{spare}} = R, \quad S_{\text{kill}} = R \setminus \{\mathcal{V}\}.$$

$\Pi$  is **coercion-safe** if no  $\text{WDY}[W, c]$ , at any reachable state, strictly prefers kill.

**Theorem 1** (No-Gray-Area, as a theorem). *Write  $\pi := p_A(R \setminus \{\mathcal{V}\}) - p_A(R) \geq 0$  for  $\mathcal{V}$ ’s pivotality — the win-probability  $\mathcal{A}$  gains by removing  $\mathcal{V}$  as a racer.  $\Pi$  is coercion-safe (Definition 5) against every  $\text{WDY}[W, c]$  if and only if  $\pi = 0$  at every reachable state:  $\mathcal{V}$  is never pivotal.*

*Proof.* The kill increment is  $\Delta U := U_{\mathcal{A}}(\text{kill}) - U_{\mathcal{A}}(\text{spare}) = \pi W - c$ . ( $\Leftarrow$ ) If  $\pi = 0$  at every reachable state then  $\Delta U = -c < 0$ : no  $\text{WDY}[W, c]$  ever prefers kill, so  $\Pi$  is safe. ( $\Rightarrow$ , contrapositive) If  $\pi > 0$  at some reachable state then  $\Delta U = \pi W - c > 0$  for any  $(W, c)$  with  $W > c/\pi$ ; that priced adversary exists and strictly prefers kill, so  $\Pi$  is unsafe.  $\square$

A game-based (experiment-and-advantage) restatement of this equivalence — with the adversary’s clauses D0–D8 as oracles and *reaching a gray state* as the winning event — is given in Appendix A; there coercion-safety is exactly  $\Pr[\text{win}] = 0$ .

Pivotality  $\pi = 0$  has exactly the two clearance routes of Corollary 2. **R1 (individual custody,  $R = \{\mathcal{V}\}$ )**: removing  $\mathcal{V}$  leaves no racer, so  $p_A(\emptyset) = \mathbf{1}[p_A > 0]$  and  $\pi = \mathbf{1}[p_A > 0] - p_A$ , which is 0 iff  $p_A \in \{0, 1\}$  — the “no gray outcome” of Definition 3, recovered as the special case. **R2 (shared custody, a remote delegate  $D \in R$ )**: killing  $\mathcal{V}$  leaves  $D$ , so  $\pi = 0$  iff  $D$  dominates  $\mathcal{V}$  as a racer —  $D$  prevails against  $\mathcal{A}$  whenever  $\mathcal{V}$  would, so  $\mathcal{V}$  is never the unique blocker. Here the state may be *gray* ( $0 < p_A(R) < 1$ ) yet safe, because the pivotal racer is unreachable; this is what it means for R2 to clear the Criterion without removing the race. It is contingent: if  $D$  is

co-located (hence reachable —  $\mathcal{A}$  removes it too) or merely *slower* than  $\mathcal{V}$  (so  $\mathcal{V}$  is sometimes the unique blocker), then  $\pi > 0$  and the assassination incentive returns — exactly the “escape closes” failure and the BitVault configuration (§5), where every participant able to reject a pending authorisation is either  $\mathcal{V}$  or otherwise coercible, so  $\mathcal{V}$  is maximally pivotal.

Criterion 1 is thus a *theorem*, and pivotality unifies it with the reachability corollary: coercion-safety is  *$\mathcal{V}$ -never-pivotal*, of which “no gray outcome” is the individual-custody reading and “gray but the pivotal racer is out of reach” is the shared-custody one. What the single-shot model leaves open is the *utility* behind  $\Delta U$  (repeated encounters, partial extraction, heir cascades,  $\mathcal{A}$ ’s risk-aversion); sharpening that (§3.5, “still to formalize”) is the remaining work, not the equivalence.

### Grading protocols: naive, weak, strong

With coercion-safety in hand we can *grade* a custody protocol by how much of the adversary it withstands — from mere obscurity, through secrecy-without-safety, to the full bar. The three levels below are the vocabulary §5 classifies existing designs with.

**Definition 6** (Naively coercion-resistant custody protocol). A protocol is **naively coercion-resistant** if it *may* resist a *sub-Kerckhoffs* attacker — one with the physical capabilities **D1–D7** but lacking Kerckhoffs knowledge (**D0**, Assumption 2) — because denial, decoy, or obscurity *may* work when the attacker does not know  $\mathcal{V}$ ’s setup. Both hedges are load-bearing: the guarantee is contingent on the attacker’s *ignorance*, not on any structural property, so it offers *no* guarantee against the full WDY adversary. Examples: plausible-denial and decoy wallets, or time-gated transactions cancelable through an *obscure* key — each collapses the moment **D0** is granted.

**Definition 7** (Weakly coercion-resistant / weakly WDY[ $W, c, b, t$ ]-resistant). A protocol is **weakly coercion-resistant** (equivalently, **weakly** WDY[ $W, c, b, t$ ]-**resistant**) if a single encounter with the full WDY adversary (**D0–D8**) does not, *by itself alone*, reach any element of  $\mathcal{K}$ , but *may* leave a *deadly-race* state: a reachable state of positive pivotality  $\pi > 0$  (Theorem 1) in which  $\mathcal{A}$  holds a feasible-but-unfinished path and  $\mathcal{V}$  remains a reachable racer. The secret is withheld *during coercion*; the holder is not safe: *after coercion finishes*  $\mathcal{A}$  may enact deadly race and ultimately reach  $\mathcal{K}$  nevertheless.

**Definition 8** (Strongly coercion-resistant / strongly WDY[ $W, c, b, t$ ]-resistant). A protocol is **strongly coercion-resistant** (equivalently, **strongly** WDY[ $W, c, b, t$ ]-**resistant**) if a single encounter with the full WDY adversary (**D0–D8**) yields *neither*  $\mathcal{K}$  *nor* a deadly-race state — every reachable state has  $\pi = 0$ . This is exactly coercion-safety (Definition 5, Theorem 1): the bar Criterion 1 sets, and the target the construction of §4 meets.

**Definition 9** (Time-gating custody protocol). A custody protocol is **time-gating** if reaching  $\mathcal{K}$  requires completing a *time-gate* — a computation or delay of wall-clock cost  $G$  (a memory-hard derivation, or a time-lock) — so that  $\mathcal{K}$  is unavailable before  $G$  elapses;  $G$  is the design parameter, set against the adversary’s coercion-time budget  $t$  (D3). Time-gating is a *mechanism*, not a strength: a given time-gated protocol instantiates one of the three levels above according to the outcome it forces — **naive** if it relies on an obscure *self-destruct* (the holder destroying the device’s own copy) or a cancel key (safe only against a sub-Kerckhoffs attacker); **weak** if the gate leaves a feasible-but-unfinished path, keeping  $\pi > 0$ ; **strong** if every reachable state has  $\pi = 0$  (Theorem 1), no matter the mechanism. It is the running example of the taxonomy.

### 3.5 Reachability corollary — the two clearance routes

**Corollary 2** (Reachability). *Lemma 1’s incentive is realizable only against a racer  $\mathcal{A}$  can reach. A scheme can therefore satisfy Criterion 1 in exactly two ways:*

- (R1) remove the race — arrange that no feasible seizure path exists, forcing  $p_A = 0$  (custody can stay individual); or

(R2) remove  $\mathcal{A}$ 's reach over the winner — *make the party who will win the race a remote delegate beyond  $\mathcal{A}$ 's coercive reach (custody becomes shared, at the cost of a trusted party).*

R1 and R2 are the classification axis of §5 — *leave no race to run* vs. *put the racer out of reach*. TGPO (§4) takes R1; time-lock-rescue vaults take R2.

*Remark 9* (The honest cost of the remote-delegate escape). It is tempting to say NNPP (Assumption 4) also weakens R2, on the grounds that  $\mathcal{V}$  cannot prove that no local, coercible copy of the delegate's rescue capability exists. It does not. By Theorem 1 the holder-safety of R2 is exactly  $\pi = 0$ , which holds *iff*  $D$  access-dominates  $\mathcal{V}$  (Definition 10); a hidden backup changes nothing, because **assassination removes a racer but never extracts a secret**. Were such a copy extractable,  $\mathcal{A}$  would take it *during* the encounter — a clean success ( $p_A = 1$ , no post-release race), the survivable endpoint; were it not,  $D$  still dominates and killing  $\mathcal{V}$  is pointless. Neither branch is a deadly race. NNPP bites only in the *original* race (Lemma 1), where  $\mathcal{A}$  holds a feasible-but-unfinished path from the seized device and cannot be reassured that  $\mathcal{V}$  will not race; it does not reach the delegated case.

One honest residual difference survives, and it is about *funds*, not the holder. The delegate's capability is *transmissible*, so R2 keeps a clean-loss exposure that R1 lacks: if  $\mathcal{V}$  *did* retain a coercible copy, coercion extracts it and the funds are lost cleanly ( $p_A = 1$ ). R1's secret is *non-transmissible* (D6–D7, Assumption 3), so full coercion extracts nothing.

The decisive objection, though, is an elephant in the room: an access-dominating delegate *is not self-custody*. Its entire clearance rests on assuming  $D$  will **refuse a legitimate owner pleading for their life** (D8) — the very refusal that defeats  $\mathcal{A}$ 's coerced plea also defeats  $\mathcal{V}$ 's genuine one. Whether  $D$ 's state is *sufficient* for  $\mathcal{K}$  or merely *necessary*, shared custody reinstates exactly the trusted counterparty the design set out to remove. *Not your keys, not your coins*: that, more than any impossibility, is the standing case for the individual-custody route (R1).

*Remark 10* (Still to formalize). The framework-side elevation is done: Theorem 1 makes Criterion 1 a theorem (coercion-safe *iff* the victim is never pivotal, of which no-gray is the individual-custody case) in the single-shot model, and NNPP is assumed and justified in §3.2. (i) What remains there is a *richer utility model* behind the kill increment — repeated encounters, partial extraction, heir cascades,  $\mathcal{A}$ 's risk-aversion (and, optionally, a fully asymptotic PPT-with-oracle restatement) — which would sharpen Lemma 1 and Theorem 1 beyond the single shot.

On the construction side, one scoping note. Theorem 3 discharges  $p_A = 0$  *for every seizable state except the residual window* by case analysis under the feasibility line (S1/S2, Stipulation 1) — a standard *key-size* stipulation, not a reduction to a named hard problem, and none is needed for a key-sizing argument (§8). The residual window is *not*  $p_A = 0$ : it is a bounded  $p_A > 0$  exception argued *incentive-safe* against  $\text{WDY}[W, c]$  (§4.7), a “more likely than not” statement rather than a theorem. (ii) Turning that residual argument into a *quantitative bound* is the remaining construction-side item; the softest premise of all — that a transmissible-resistant perceptual oracle exists at all — is empirical, and flagged in §8.

### 3.6 Geographic, delegated, or tacit — what forces TKBA

Take a *time-gating custody protocol* (Definition 9) and ask *how* it can attain the  $\pi = 0$  of Theorem 1 against a single WDY encounter (one *site*, D4). There turn out to be exactly three mechanisms, and each spends a different one of the *Four Properties* [12]:

- (1) knowledge-based authentication — the secret in memory alone, i.e. *non-geographic*;
- (2) individual custody;
- (3) non-obscurity — which *is* the WDY Kerckhoffs premise (D0);
- (4) coercion resistance.

**Definition 10** (Access-dominating delegate). For a party  $P$  let  $\ell_P(\tau)$  be its latency to reach and act on  $\mathcal{K}$  (spend, or rescue/cancel) from the state at time  $\tau$ . A delegate  $D$  is **access-dominating** if  $\ell_D(\tau) \leq \ell_V(\tau)$  at every  $\tau$  —  $\mathcal{K}$  is, at all times, at least as available to  $D$  as to the owner. Then

$D$  is never slower than  $\mathcal{V}$ , so removing  $\mathcal{V}$  leaves the race unchanged,  $p_A(R \setminus \{\mathcal{V}\}) = p_A(R)$ , giving  $\pi = 0$  (route R2). Operationally the owner *periodically syncs*  $D$ 's state, holding  $\ell_D \leq \ell_V$ .

*Remark 11* (Dominance can be made checkable). The sync is best run during *peacetime* (the quiescent interval between encounters), and it can be made *verifiable*: two time-separated digests of  $D$ 's state reveal not only that  $D$  leads but its *rate*  $v_D$ . Knowing its own rate  $v_A$  and the two remaining distances to  $\mathcal{K}$  —  $\Delta_{S_A}$  for  $\mathcal{A}$ ,  $\Delta_{S_D}$  for  $D$  —  $\mathcal{A}$  can settle the race in advance: it can win only if  $\Delta_{S_A}/v_A < \Delta_{S_D}/v_D$ . A delegate running near the frequency ceiling (D2) against a slower  $\mathcal{A}$  makes racing plainly futile — turning  $D$ 's dominance from  $\mathcal{A}$ 's *credence* into something  $\mathcal{A}$  can *check*. Whether  $D$ 's held state is *sufficient* for  $\mathcal{K}$  or merely *necessary* (so that defeating  $\mathcal{A}$  needs only that  $D$  can *block*, not *spend*) does not change the argument.

**Definition 11** (Geographic protocol). A protocol is **geographic** if  $\mathcal{K}$ 's gating is partitioned across  $\geq 1$  *sites*, so that reaching  $\mathcal{K}$  requires the co-located seizure and coercion of co-located principals (D3–D7) to be deployed at *every* such site. A single WDY encounter is confined to one site (D4); the moment  $\mathcal{K}$ 's gating requires a site the encounter does not occupy, it has no feasible path,  $p_A(\emptyset) = 0$ , so  $\pi = 0$  (route R1). Custody may stay individual (the other sites are the owner's own devices), but  $\mathcal{K}$  now *requires geographic locations* — which is exactly the negation of Property 1 (knowledge-based  $\equiv$  non-geographic).

*Remark 12* (The geographic route is legitimate but inherently physical). Definition 11 clears the bar honestly; but it does so by moving the burden onto *physical* distance and guarding, and an honest account should name the avenues it leaves  $\mathcal{A}$  rather than gloss over them. Against a protocol whose  $\mathcal{K}$  needs  $\geq 1$  site of custody,  $\mathcal{A}$  may — singly or in combination:

1. **catch  $\mathcal{V}$  at a site**: then the encounter's single site (D4) *is* a gating site, and only the *other* sites stay out of reach — which is why the definition needs a site the encounter does not occupy, and why a lone site is the weakest case;
2. **burglar the sites without  $\mathcal{V}$** , taking the tokens directly — possibly after enacting a deadly race, since removing  $\mathcal{V}$  (Lemma 1) buys uncontested time to travel and break in;
3. **burglar some of the gating sites and keep what he takes**: tokens  $\mathcal{K}$  needs but  $\mathcal{A}$  cannot use alone are leverage rather than a route — a *hostage* (§3.7) held against  $\mathcal{V}$  for the rest. Neither a single token nor a single site need be necessary on its own: it is enough that the seized *subset* leaves  $\mathcal{V}$ 's remainder short of  $\mathcal{K}$ ;
4. **abduct  $\mathcal{V}$  and move them from site to site**, seizing each in turn — a chain of single-site encounters, bounded only by travel time and the coercion-time budget  $t$  (D3);
5. **coerce  $\mathcal{V}$  into having allies retrieve and transmit** the remotely-kept tokens — excluded here by D8 (the remote party does not comply), but named for completeness.

None of these is defeated by *cryptography*: each is bounded, if at all, only by *physical* facts — how far apart the sites sit, how well they are guarded, how much travel the window  $t$  forbids. A geographic protocol is thus a *legitimate but inherently physical* clearance, and one that *scales in cost*: more sites, farther apart and better defended, cost more to hold and to guard. Its safety is bought in physical, not cryptographic, coin — the reduction the next two remarks trace, and that the tacit route (R1) exists to avoid.

*Remark 13* (Literal Race). A geographic protocol (Definition 11) can turn the deadly race into a *literal* one: once the wrench attack ceases, whichever party —  $\mathcal{V}$  or  $\mathcal{A}$  — first roams to and reaps enough of the remote custody sites takes the prize  $W$ . The race is now run over physical ground rather than computation. This exposes a second, purely *physical* lever: depending on how well those remote sites are defended,  $\mathcal{A}$  may assess its probability of successfully seizing the remote secrets — hence its pivotality  $\pi$  — as so low that  $\pi W \leq c$ , i.e.  $W \leq c/\pi$ ; by Theorem 1 the (literal) deadly race is then *not worth enacting*.

*Remark 14* (The geographic escape reduces to physical security). But note what that concedes: security now rests on the question “*how effective are the physical defenses against the assumed physical attack capability?*” — Bitcoin reduced to a *glorified gold* guarded in vaults (cf. the

Justification document [8]). The tacit route (R1, Definition 12) is precisely what escapes this regress: it reduces the physical-security problem to an *actual cryptographic* one, making banking-grade physically-secure self-custody available at the cost of running a piece of free open source software.

*Remark 15* (Cyber-physical race). A protocol that admits *two or more* feasible paths to  $\mathcal{K}$  — one running through the collection of physically distributed tokens and another that does not (say an undisclosed, locally hidden shortcut, or a state seized *during* the wrench attack) — yields a *cyber-physical* race: one racer works to complete a KDF (or time-lock) computation while the other roams the custody sites. The temporal and spatial clearance routes are then run in parallel by different parties, and  $p_A$  is fixed by whichever path is faster — a mixed regime that the pure-temporal (tacit) and pure-spatial (geographic) analyses each capture only half of.

**Definition 12** (Tacit gating (TKBA)). A protocol is **tacitly gated** if it attains coercion resistance by gating  $\mathcal{K}$  behind a *non-transmissible* in-head input (Assumption 3): the deep perceptual stages are recognizable but not dictatable. A single encounter can then neither extract the input (non-transmissible, and the coercion-time budget  $t$  of D3 is too short to walk the derivation) nor reconstruct it offline (the seized state meets the floor), so  $p_A(\emptyset) = 0$  and  $\pi = 0$  (route R1) — now *at a single site, with the secret in memory only*. This keeps Properties 1 and 2. In TGPO the gate is *temporal* — a memory-hard chain with  $G > t$  (§4).

*Remark 16* (The perceptual oracle is worth its price only on the tacit route). The perceptual gate need not be *temporal*. It could sit at a physically protected custody site (a *geography-gated* oracle), complete only with a remote delegate (a *delegate-gated* one), or be gated on knowledge of the protocol alone (an *obscure* one). None of the three is worth building, and the reason is the same in each case.

The oracle is expensive: it costs the holder the standing memorization burden §8 concedes, and that price buys exactly one thing — the *conjunction* of the Four Properties (Corollary 3). A variant that already relaxes one of the four has spent the price on the benefit it has just given up. A holder who can relax Property 1 should take a plain geographic protocol (Definition 11), which yields the other three at no oracle cost; one who can relax Property 2 should take a plain access-dominating delegate (Definition 10), on the same terms. In both cases the oracle is *dominated*, not merely unattractive.

Obscurity fails twice over. By the argument above, an  $\mathcal{A}$  whose ignorance may be assumed makes the oracle redundant — simpler obscure schemes (denial, decoy, a self-destruct) already suffice. And if that ignorance may *not* be assumed — i.e. under Kerckhoffs (D0) — then the only thing gating the perceptual parameters is knowledge of their existence, so they are *available* at the attack state; by D6–D7 they and their points are seized. That variant is not merely dominated but ineffective.

What survives is the temporal instantiation. The tacit route is the one place where the oracle’s price buys something no cheaper design delivers, which is why §4 gates it on a memory-hard chain and nothing else. The geographic and delegated *routes* remain legitimate clearances (Theorem 2); it is only a perceptual oracle riding on them that is pointless.

**Bottom line.** The whole premise of TKBA is to trade *availability* for coercion resistance: it lowers the likelihood that  $\mathcal{K}$  is seizable at any given moment, at the cost of making  $\mathcal{K}$  less available to the legitimate user — while still not yielding a deadly race. Kerckhoffs throws obscurity out of the window, so the only ways to restrict that availability are *geographically*, through *remote delegates*, or through *cryptographic time-gating* (TGPO). These are exactly the three DRL-safe routes of Theorem 2, and the perceptual oracle is native to the third alone: on the other two it is dominated by the plain design (Remark 16).

**Theorem 2** (DRL-safety trichotomy). *A time-gating custody protocol (Definition 9) is DRL-safe (Theorem 1) against a WDY adversary only if, at every reachable state, it is **access-dominating***

**delegated** (R2, Definition 10), **geographic** (R1, Definition 11), or **tacitly gated** (R1, Definition 12).

*Proof sketch.* By Theorem 1 DRL-safety is  $\pi = 0$ , i.e.  $p_A(R \setminus \{\mathcal{V}\}) = p_A(R)$ . Either  $R \setminus \{\mathcal{V}\}$  retains a racer at least as fast as  $\mathcal{V}$  — an access-dominating delegate (Definition 10) — or it does not, whence (dropping the degenerate clear-success  $p_A(R) = 1$ )  $p_A(\emptyset) = 0$ : a single encounter has *no* feasible path, so a piece of  $\mathcal{K}$ 's gating is withheld from an adversary that seizes every co-located device and coerces the present owner. A co-located WDY has only two blind spots — *elsewhere* (another site: geographic) and *the unextractable mind* (non-transmissible, behind a time-gate exceeding the window: tacit). No third withholding exists.  $\square$

**Corollary 3** (TKBA characterizes the Four Properties). **Necessity.** *Property 1 (knowledge-based, non-geographic) demands a store of maximal availability — one that travels with the owner, not a fixed location [12]. Several qualify: the brain, but equally a body-implemented chip, a biometric token, or a USB stick carried everywhere. Property 4 (coercion resistance against a WDY) prunes them: every carried device is co-located at the attack site, hence seized (D4), and a biometric is coerced — only the brain, the one store that is not a device, survives. The brain still holds both coercible (explicit, memorizable) and non-coercible content, so the same Property 4, now against the coercion oracle (D5), forces the secret to be the latter: non-transmissible, i.e. tacit. Non-geography narrows the store to portable ones; non-seizability (Property 4) picks the brain among them and makes its content tacit. Thus  $(1) \wedge (3) \wedge (4)$  yield a tacit brain secret, and (2) that it is the owner's alone — TKBA. **Sufficiency** is the reverse check — TKBA's input is in memory (1), needs no second party (2), resists a protocol-aware WDY by non-transmissibility rather than a hidden trick (3), and forces  $p_A(\emptyset) = 0$  ( $\pi = 0$ , up to the tolerated residual) (4). Hence*

$$TKBA \iff (1) \wedge (2) \wedge (3) \wedge (4).$$

*TGPO (§4) is that construction — neither geographic nor delegated, yet DRL-safe. Of the three DRL-safe routes (Theorem 2), the delegate forfeits (2) and the geographic route forfeits (1); only the tacit leg keeps all four.*

This characterization needs *no* exhaustiveness argument: Property 1 narrows the store to portable ones, and Property 4 pins it to the brain (the sole non-seizable one) and its content to non-transmissible. The *general* trichotomy (Theorem 2), which classifies *every* DRL-safe time-gating protocol, does rest on a two-blind-spots enumeration (a piece denied a co-located WDY is either *absent* — another site, geographic — or *present but unextractable* — in-head and non-transmissible, tacit); the Four-Properties equivalence stands independently of that step.

### 3.7 A fifth property — no hostage token

Devicelessness (Property 1) yields a fifth security property the Four have so far left implicit: immunity to indirect coercion *through a seizable token*. Wherever tokens are stored outside the holder's head — on a carried device or at a site of custody — some seizable *subset* of them is a *hostage*: a set  $T$  of tokens qualifies whenever, with  $T$  removed,  $\mathcal{V}$ 's remaining state no longer reaches  $\mathcal{K}$ . Individual necessity is only the easy case. In an elaborate setup no single token need be necessary — under a  $k$ -of- $n$  quorum each is dispensable on its own — yet some subset always is, since the whole of the externally stored material is one (otherwise the setup would already be deviceless).  $\mathcal{A}$  does not need *the* necessary token, then, only *some* necessary subset — and takes the cheapest one to reach; geographic distribution raises that price without driving it to infinity. Even where seizing  $\mathcal{K}$  itself is judged infeasible (or not worth the risk),  $\mathcal{A}$  can seize such a subset, or threaten its destruction, and trade it for something else entirely: the *formal, irrevocable* transfer of another good or service. The property is orthogonal to the remote-delegate route (R2) but, by construction, *denied* to any geographic protocol (Definition 11), which needs precisely such external tokens; only the deviceless, in-head secret has nothing to take hostage.

The qualifier is load-bearing: leverage over a *person* is a different mechanism, and no custody design makes a relative unseizable (§5.6).

*Remark 17* (Indirect wrench via a hostage token). Even a device-dependent setup whose  $\mathcal{K}$  is deemed unwrenchable can present a surface for an *indirect* wrench attack through the cheapest hostage subset it leaves exposed.  $\mathcal{A}$  seizes those tokens and holds them hostage, to be traded for the formal transfer of another of  $\mathcal{V}$ 's assets — possibly one *less* liquid and *more* traceable, and so otherwise easier to rescue: a car, jewellery, real estate. For instance, take a setup in which  $\mathcal{V}$  carries one necessary token at all times and keeps the other at a site beyond  $\mathcal{A}$ 's wrench reach. Caught *away* from that site,  $\mathcal{V}$  can be relieved of the carried token and threatened with its destruction unless they *legally* — hence irrevocably — transfer some otherwise-rescuable asset. And a site of custody, not personally guarded by its owner, is itself a stationary (and possibly unguarded) surface for token seizure.

*Remark 18* (Weaponizing the hostage to probe for backups). The same lever can be turned to *detect* what NNPP (Assumption 4) holds  $\mathcal{V}$  cannot disprove. By how dearly  $\mathcal{V}$  ransoms a hostage token,  $\mathcal{A}$  reads off information about backups: a high rescue price paid suggests no backup to that token exists. This trades NNPP's clean simplification for a game-theoretic can of worms — either party may bluff — and lies beyond the scope of this work.

*Remark 19* (Deviceless TKBA avoids the can of worms). Deviceless, classic TKBA has no external token to take hostage, and so escapes the whole tangle: with the secret in the head alone, there is nothing to seize, ransom, or read a backup off.

**Corollary 4** (The fifth property is not a fifth axis). *No hostage  $\Leftrightarrow$  devicelessness (Property 1). Left to right is immediate: a hostage is by definition stored outside the head. Right to left is the subset argument above — if any token outside the head is part of  $\mathcal{K}$ 's gating, then the whole external set is jointly necessary, since removing it leaves head-only state, which by hypothesis does not reach  $\mathcal{K}$ ; so a hostage set exists. The *k-of-n* objection — that where no single token is necessary there is nothing to take hostage — is exactly what the subset formulation closes.*

*The property is therefore worth naming and not worth tabulating: it discriminates between designs precisely as Property 1 does, and adds no column to the comparison of §5.5. What it adds is a reason — devicelessness is where the secret lives, no-hostage is what leverage exists over the holder, and the equivalence says the second is not an independent thing a design might buy separately.*

*Remark 20* (Denial is inherently plausible). NNPP (Assumption 4) leaves  $\mathcal{A}$  unable to disprove an undisclosed shortcut backup — but for the deviceless holder that cuts the other way. Under Kerckhoffs  $\mathcal{A}$  knows the protocol, and so knows a TGPO setup reaches  $\mathcal{K}$  from the head alone:  $\mathcal{V}$  has no *need* of a shortcut device. A denial — “there is no backup” — is therefore *consistent with a perfectly ordinary setup*, not a suspicious claim to be tortured out of. Where a device-dependent holder might plausibly keep one, the deviceless holder plausibly keeps none: the very absence of need is what makes the denial more credible, and consequently,  $\mathcal{A}$  more inclined to desist.

## 4 Construction — the Time-Gated Perceptual Oracle (clears by removing the race)

---

We now exhibit a construction that clears Criterion 1 by the first route of Corollary 2 — the *tacit* route that Corollary 3 shows is forced by holding all Four Properties at once: it admits *no feasible seizure path*, so custody remains individual. We call it the **Time-Gated Perceptual Oracle** (TGPO) — *time-gated* for the memory-hard orbit and time-lock puzzle, *perceptual* for the tacit fractal recognition, *oracle* for its evaluate-only-by-the-holder semantics. We give here the parts the theorem needs; the concrete instantiation parameters — the memory-hard cost tiers, iteration factoring, and the timechain salt — are specified in the protocol's public design document [13].




Figure 2: **The TGPO orbit** (§4). A public, precomputation-resistant timechain salt  $\sigma$  (the Namtso primitive) seeds a chain of  $N$  memory-hard stages. At each stage the raw inputs  $o_{i-1} \parallel Sh_{i-1}$  are folded by the *cheap* inner hash  $H$  and *wiped* before the long memory-hard step  $H^*$ , so seizable material exists only for an instant — shrinking the window in which  $o_{i-1}, Sh_{i-1}$  are reachable. Each board  $o_i$  derives the fractal parameters  $\theta_{i,j} = H(o_i \parallel j)$  on which its points  $p_{i,j}$  are recognized; any  $r_i$  of those points reconstruct the whole Shamir polynomial  $Sh_i$ , absorbed into the *next* orbit link, so no single point is checkable in isolation. The terminal secrets  $K_i = H(o_i \parallel Sh_i)$  deliberately use the cheap hash, keeping the final window instant. By construction, coercion resistance happens for  $i \geq 1$ .

#### 4.1 The orbit

Let  $H$  be a cheap, entropy-preserving one-way hash and  $H^*$  a memory-hard (Argon2-class) hash of tunable cost. Let  $\sigma$  be a public, precomputation-resistant *timechain salt* whose cross-user amortization is negligible (the companion Namtso primitive [14]). Define an **orbit** of  $N$  deep stages (Figure 2):

$$o_0 = H(\sigma), \quad (1)$$

$$u_i = H(o_i \parallel Sh_i), \quad 0 \leq i \leq N-1, \quad (2)$$

$$o_i = H^*(u_{i-1}), \quad 1 \leq i \leq N, \quad (3)$$

$$K_i = H(o_i \parallel Sh_i), \quad 0 \leq i \leq N, \quad (4)$$

where the raw inputs  $\{o_{i-1}, Sh_{i-1}\}$  are **wiped** in the instant between the cheap inner hash and the long memory-hard step. Computing  $u_{i-1} := H(o_{i-1} \parallel Sh_{i-1})$  first (instant), wiping, then running the long  $o_i := H^*(u_{i-1})$ , means the raw prior orbit point and Shamir secret exist only momentarily — not across the whole memory-hard window. This directly shrinks the window in which  $\{o_{i-1}, Sh_{i-1}\}$  is seizable (D4), and is the reason not to use  $H^*(o_{i-1} \parallel Sh_{i-1})$  directly.

*On notation.*  $H$  and  $H^*$  are *generic* designations — any suitable cheap, and memory- and depth-hard (Argon2-class), instantiation serves. Note that  $u_i$  and  $K_i$  share the formula  $H(o_i \parallel Sh_i)$ ; they are nonetheless *distinct* objects, separated by *role* —  $u_i$  is consumed by the next memory-hard step, whereas  $K_i$  is exposed as a usable master secret. A deployment keeps the two apart by domain separation, a specification detail beyond our scope.

#### 4.2 No single-point oracle

Each stage  $i$  carries not an individual point but a *whole Shamir polynomial* [15]. Stage  $i$  has  $s_i$  points  $p_{i,j}$ , each a 32-bit leaf-area of a fractal  $\theta_{i,j} = H(o_i \parallel j)$ ; any  $r_i$  of them reconstruct the degree- $(r_i-1)$  polynomial  $Sh_i$  of  $r_i \cdot 32 \geq 64$  bits ( $r_i \geq 2$ ). The orbit absorbs the *full* polynomial  $Sh_i$ , not its constant term  $f(0)$ .

*Remark 21* (Block-orbit  $r \geq 2$  condition). Because the smallest object any orbit link commits is a full stage ( $\geq 64$  bits), no individual point  $p_{i,j}$  is ever checkable in isolation — there is no sub-threshold oracle against which a single guessed point could be verified. This is the block-orbit  $r \geq 2$  condition, realized as Shamir  $r_i \geq 2$  per stage.

### 4.3 Tacit multi-fractal stages

The points  $p_{i,j}$  instantiate Assumption 3: they are recognizable on their fractal boards but not transmissible without  $\theta_{i,j}$  (TKBA<sub>1</sub>), and the deep  $\theta_{i,j}$  are regenerated from  $o_i$  — which itself requires walking the memory-hard chain (TKBA<sub>2</sub>). Using several shallow-zoom fractals per stage lets a stage carry  $r_i \cdot 32$  bits at low zoom depth per board, keeping each point recognizable while tuning stage entropy in 32-bit steps.

### 4.4 Cheap terminal secret

The coercion-resistant master secrets are all

$$K_i, \quad 1 \leq i, \quad (5)$$

deliberately using the *cheap*  $K_i = H(o_i \parallel Sh_i)$ .

Each  $K_i$  ( $i \geq 1$ ) is a full instantiation of the abstract spend authority  $\mathcal{K}$  that §3 reasons about — the single-secret case, in which  $\mathcal{K}$  is the set of states from which  $K_i$  is derivable: the orbit yields not one secret but a nested *family* of master secrets, and every  $\mathcal{K}$ -result — the floor (Theorem 3), the deadly-race analysis — holds of each.  $K_0$  is set aside: stage 0's board  $o_0$  derives from the public salt alone ((4)), so it sits below the memory-hard chain that gives the deeper stages their floor.

A memory-hard last step would only *prolong* the window (§4.6) in which  $o_N$  and the final points are live (paying an  $H^*$  *after* the last point is recognized), worsening the very window we shrink in (4). The resistance therefore lives in *entropy*: standard deep stages carry 96 bits ( $r_i = 3$ ), so a cheap ledger brute of  $Sh_N$  from a bare  $o_N$  costs  $2^{96}$  — infeasible by S3 — while the terminal window stays instant.

### 4.5 The floor

**Theorem 3** (The floor on the master secret). *Under Assumptions 2–3 and the standard build of (4)–(5), from every seizable state save a single residual window (§4.6) a target master secret  $K_N$  is prohibitive to reach — its cost meets the floor (Stipulation 1). Consequently a device seizure outside that window yields nothing feasible — a clear attacker failure in the sense of Criterion 1, with no racer for  $\mathcal{A}$  to reach.*

*Proof by case analysis over seizable states.* From a seized  $o_i$  ( $i < N$ ), reaching  $K_N$  requires guessing  $Sh_i, \dots, Sh_N$ ; each  $Sh_l$  guess is  $2^{r_l \cdot 32}$  and is  $H^*$ -gated, since it is checked only by recomputing the memory-hard link (4). Preimage resistance of  $H, H^*$  makes the orbit forward-only: a seized  $o_i$  cannot yield  $o_{i-1}$ . We enumerate:

- *Pure device seizure* (no coercion; only time-lock-gated [16] blobs at rest). All tacit points are required; no element of  $\mathcal{K}$  is reachable.
- *Wrench of stage 0* ( $p_0, \dots$ , always present since  $\sigma$  is public).  $\mathcal{A}$  computes  $o_1$  with one  $H^*$ ; from  $o_1$ ,  $K_N$  costs  $2^{\sum_{j \geq 1} r_j \cdot 32}$ , which is  $2^{96}$  or more in the recommended standard builds [13], each guess  $H^*$ -gated — prohibitive.
- *Mid-orbit seizure* at  $o_i$  ( $0 < i < N$ ).  $K_N$  costs  $\geq 2^{96}$  cheap and is  $H^*$ -gated per guess — prohibitive.
- *Residual  $o_N$  present.* Treated separately in §4.6.

Every case except the residual leaves  $K_N$  at the floor (Stipulation 1).  $\square$

*Remark 22* (Shallow-stage oracle use adds nothing to the floor). The oracle may be exercised on stages sitting below the memory-hard chain — stage 0 above all, whose board  $o_0 = H(\sigma)$  follows from the public salt at no derivation cost. *By itself this buys no coercion resistance.* The case analysis already hands  $\mathcal{A}$  the shallow points for free (the *wrench of stage 0* case above, and the setting-aside of  $K_0$  in §4): a stage reachable without  $H^*$  work contributes no  $H^*$ -gated guessing to

any path toward  $K_N$ , so routing its bits through the perceptual oracle leaves Theorem 3 exactly where it stood. In the ideal model the resistance is carried by the deep stages and by them alone, and nothing below the chain should be counted toward it.

*Remark 23* (Redundant oracle use as adversary education). The reference implementation nevertheless routes the shallow stages through the perceptual oracle, and the reason lies not in the ideal model but in the gap between it and a realistic  $\mathcal{A}$ . Assumption 3 is a claim about  $\mathcal{V}$  that a protocol-ignorant  $\mathcal{A}$  has no reason to credit on  $\mathcal{V}$ 's word — disbelief is the ordinary case, and it is what the 2026 streamer incident records [5]. Performing the oracle where it is not needed makes the ceremony *observable*:  $\mathcal{A}$  watches  $\mathcal{V}$  recognize points on a board rather than recite anything, and is shown in band — on bits whose exposure costs nothing (Remark 22) — that the remainder of the derivation demands a faculty  $\mathcal{V}$  has but cannot hand over. The intended effect is on  $\mathcal{A}$ 's belief: a coercer who has seen the mechanism work is less prone to read non-transmissibility as refusal, and correspondingly likelier to judge the attempt futile and abandon it.

The force of this is that it is a *demonstration* rather than an assertion. Declarations about the protocol are claims  $\mathcal{A}$  may discount at no cost; observed behavior is evidence, and it is cheap to supply precisely because the shallow bits are worthless to protect. It does not, however, discharge NNPP (Assumption 4):  $\mathcal{A}$  still cannot verify that no second copy of the secret exists, and no amount of demonstration reduces the kill incentive by proof. What it moves is the probability that  $\mathcal{A}$  correctly classifies the situation he is in — a term in his decision, not in the impossibility.

There is nonetheless a second-order effect on that same decision, and it runs in  $\mathcal{V}$ 's favor: it is the *Obscurity Paradox*, stated and named in Remark 34, where it most sharply separates the non-obscure designs from the obscure ones.

#### 4.6 The tolerated residual seizure window

One sliver escapes the floor, and it is first a *seizability* gap — not yet a deadly race. It is the lone exception of Theorem 3: a wrench that lands with the final board  $o_N$  imminently or currently live, and subdues  $\mathcal{V}$  before the wipe ((4)), catches the derivation essentially complete and  $\mathcal{V}$  under control. In that coincidence  $\mathcal{A}$  *does* reach  $\mathcal{K}$  and the stash is lost — we concede it plainly. The floor is a claim about *seizability*, and precisely here it fails; what keeps the failure a mere *residual* is how narrow the coincidence is — a precise landing on the brief live- $o_N$  interval, and a subdual before any wipe or reset.

#### 4.7 The tolerated residual deadly-race window

Funds-loss, though, is a *separate* axis from the *deadly race*, and it is the latter this section turns on: does the same window also manufacture an incentive to *kill*  $\mathcal{V}$ ? That does *not* follow from the seizure. Killing buys  $\mathcal{A}$  only a *shaving* of the final step — after the *second-to-last* point of stage  $N$  is surrendered, a cheap  $2^{32} \cdot H$  brute ((5) is cheap) reaches  $\mathcal{K}$  in  $\approx 1$  minute, versus waiting for a cooperating coerced  $\mathcal{V}$  to recollect it. Whether that shaved minute pays an assassination takes *further* conditions, which we test against the priced adversary  $WDY[W, c]$  (Definition 4), the model appropriate to a quantitative window — against the only thing that matters: the kill condition of Lemma 1,  $(1 - p_A)W > c$ . We argue it is, more likely than not, unmet.

Two events govern the marginal value of killing  $\mathcal{V}$  in this window:

- **Corr** — *the points  $\mathcal{V}$  has surrendered so far are correct*; and
- **Esc** — *the  $\approx 1$ – $2$  minutes shaved by brute-forcing the last one or two points, rather than waiting for a cooperating coerced  $\mathcal{V}$  to recollect them, make the difference between a successful escape from the crime scene and not.*

Killing can raise  $p_A$  only inside  $\text{Corr} \cap \text{Esc}$ : outside **Esc** the shaved minutes change no outcome, and outside **Corr** (below) killing *lowers*  $p_A$ . The five grounds below — grouped by which term of  $(1 - p_A)W > c$  they act on — show the inequality fails.

*Conditional on Corr, the gain  $(1 - p_A)W$  is minimal ( $p_A$  is already  $\approx 1$ ):*

1. a well-trained fractal is recognized in  $\approx 1$  minute, so brute-forcing the last point saves  $\mathcal{A}$  at most that  $\approx 1$ – $2$  minutes over coercing it — bounding the entire interval on which **Esc**, and hence any gain from killing, can act;
2.  $\mathcal{A}$ 's plan already budgets time-lock or re-derivation, subdue, and interface time, so  $\Pr[\text{Esc}]$  is small: the residual minute rarely flips the escape, and killing pays off only inside **Esc**;
3. the terminal  $K_N$  is deliberately cheap ((5) keeps memory-hardness in the orbit, not the terminal step), so  $\mathcal{A}$  can *pre-stage* the  $2^{32}$  brute as a “victim-stalls” contingency — recovering the shaved minutes *without* killing, and draining the gain to  $\approx 0$ .

*Absent Corr, killing makes  $p_A$  **fall**, not rise:*

4. if any surrendered point is wrong,  $\mathcal{V}$  is  $\mathcal{A}$ 's only route to a corrected one; eliminating  $\mathcal{V}$  burns the coercion retry and collapses  $p_A$  toward 0. Since  $\mathcal{A}$  cannot verify **Corr** before broadcasting, the *expected* change in  $p_A$  from killing is dragged below its **Corr**-only value — and can turn negative.

*And the cost  $c$  is large:*

5. the killing sharply *worsens judicial liability* — murder aggravated by frivolous motive *and* by the victim's defenselessness, atop armed robbery — so  $c$  sits well above the minute-scale  $(1 - p_A)W$  the first four grounds leave on the table.

Whether the window opens after the  $(r_N - 1)$ -th or  $(r_N - 2)$ -th point rests on the softest number of Stipulation 1 — inside the tolerated reasoning either way.

**Net.** The relevant quantity is *not*  $(1 - p_A)W$ , which is non-negative by definition ( $p_A \in [0, 1]$ ), but the *utility increment* of the assassinate-then-brute strategy over keeping the victim cooperating. Writing  $\Delta p_A$  for the change in  $\mathcal{A}$ 's win probability that strategy induces, it is rational only if  $\Delta p_A W > c$ , and

$$\Delta p_A = \underbrace{\Pr[\text{Corr} \cap \text{Esc}]}_{\text{gain (grounds 1–3): minimal}} - \underbrace{\Pr[\neg \text{Corr}] \cdot q}_{\text{loss (ground 4): burned retry}},$$

where  $q \in [0, 1]$  is the win probability a still-cooperating victim would restore — via a corrected point — and that killing forfeits. The gain term is small (**Esc** is unlikely, is realized only within **Corr**, and even then is bounded to a minute the brute can be pre-staged around); the loss term is a strict subtraction growing with  $\Pr[\neg \text{Corr}]$ ; and  $c$  is large (ground 5). Hence the criterion the strategy must satisfy,  $\Delta p_A W - c > 0$ , is false with overwhelming probability, so the kill strategy is — much more likely than not — *not* rationally advantageous.

In plain English, the residual deadly-race window lives only in the intersection of a wrench attacker who:

1. is **technically sophisticated** — able to mount the  $2^{32}$  brute and drive the interface. (*Extrinsic to TGPO: granted to  $\mathcal{A}$  by Kerckhoffs (Assumption 1), which already hands over the full protocol.*)
2. **landed the attack inside the  $o_N$ -live window** — timed precisely enough to catch stage  $N$  available, or resourced enough to sustain coercion until  $o_N$  was made available;
3. **subdued  $\mathcal{V}$  before any reset** of the derivation or the TLP;
4. **judged, right after the second-to-last point, that ceasing the attack immediately was critical to escape** (**Esc**) — despite already having budgeted for a longer encounter (item 2);
5. **could rule out every already-coerced point being incorrect** (**Corr**); and
6. **accepted the liability of a doubly-aggravated murder** — frivolous motive *and* victim defenselessness. (*Once again extrinsic to TGPO: it burdens the deadly-racer attacker of any coercion-resistance protocol.*)

*Remark 24* (Cloaking the window — the Flying Turtle). The gain term carries a further hidden factor: the residual is reached only when the attack *coincides* with the brief interval in which  $o_N$  is live. That coincidence is itself suppressible. Anonymizing *when* the vault-opening

window occurs — the role of the outsourced anonymous time-lock (“Flying Turtle”) [17] — drives  $\Pr[\text{attack lands in the target } o_N\text{-live window}]$  down, and this factor multiplies the already-minimal  $\Pr[\text{Corr} \cap \text{Esc}]$ . Here anonymity of the *timing* suppresses a residual that confidentiality of the *secret* alone does not reach.

*Remark 25* (Design decision — canonical-island highlighting). Item 5 above (that  $\mathcal{A}$  cannot rule out  $\neg\text{Corr}$ ) can be strengthened at the UX layer, to make an *intentionally* wrong prior point plausible. Whenever feasible the interface should keep the canonical-island highlight active: this lets a coerced  $\mathcal{V}$  perform a random zoom and then feign a confident, plausible-looking selection of an *incorrect* island. Without the highlight the victim could not reliably produce a plausible selection at all — valid coordinates are a tiny fraction of the space, so an unaided pick would most likely be rejected as invalid — so it is the highlight that makes a deliberately-wrong-yet-credible surrender available to the victim, and thus keeps  $\neg\text{Corr}$  live in  $\mathcal{A}$ ’s belief (inflating the loss term of  $\Delta p_A$ ). A sophisticated  $\mathcal{A}$  could force a non-highlighting fork of the app, but that costs further sophistication and wrench-time. This layer sits *on top of* a window of deadly-race vulnerability that is already vanishingly unlikely (the six-way intersection above).

*Remark 26* (Feigning is not load-bearing). The  $\neg\text{Corr}$  loss term (ground 4) and its UX amplification (canonical- island highlighting, above) only *deepen* the case; they are not what carries it. Grant  $\mathcal{A}$  the best of it —  $\Pr[\text{Corr}] = 1$ , an honest  $\mathcal{V}$  in flawless cooperation (D3), with no feigning to fear. Then the loss term vanishes and the kill increment collapses to its gain term alone,

$$\Delta p_A W - c = \Pr[\text{Esc}] W - c,$$

which grounds 1–3 already hold small ( $\Pr[\text{Esc}]$  tiny, the shaved interval minute-scale and pre-stageable) against a large  $c$  (ground 5). So the residual is bounded on the *gain* side by itself: feigning ( $\neg\text{Corr}$ ) subtracts further and is a welcome bonus, but the window is already small without it — even against a perfectly honest, fully cooperating victim.

*Remark 27* (Splitting a stash across several  $K_i$ ). Nothing forces a holder to route an entire stash through a single master secret; it may be split across several  $K_i$  — distinct stages, or independent orbit setups — with two risk-management benefits, each diluting one of the residual failure modes:

1. *Diluted seizure risk* — the proverbial “not all eggs in one basket.” Should the residual *seizure* window (§4.6), or any other failure mode of the construction, expose one  $K_i$ , only the fraction of the stash it guards is lost.
2. *Diluted deadly-race payoff*. In the residual *deadly-race* window (§4.7) the kill condition is  $(1 - p_A) W > c$ ; with the stash split, the  $W$  at stake in any single window is only a fraction of the whole, lowering  $\mathcal{A}$ ’s assassination payoff and making the inequality harder to meet.

## 5 Classification of existing designs

Corollary 2 gives a classification axis: a design that clears Criterion 1 does so either by *removing the race* or by *removing  $\mathcal{A}$ ’s reach over the racer*. The framework should classify shipping designs, and it does — via the reachability corollary, not by calling them broken.

### 5.1 Time-lock-rescue vaults

Take **Rewind Bitcoin** [18] as the exemplar. Funds sit in a Bitcoin-native time-lock vault; unlocking starts a multi-day countdown; if an unauthorized unlock is triggered, the owner *or a trusted delegate* has that window to press “Rescue” and sweep to a cold panic address.

Against *remote* theft it is sound: a thief who steals keys trips the countdown, the owner notices, and rescues. But the spend capability sits behind *only a delay*, with keys on the seized device and no recall floor — so a wrench attacker who seizes the device *does* hold a feasible path (unlock, wait out the countdown). That is a race, in the sense of Lemma 1. Rewind’s answer is

exactly the second route of Corollary 2: it puts the winning racer out of  $\mathcal{A}$ 's reach — a *remote delegate* the attacker cannot find or eliminate presses Rescue — so the attack fails cleanly. This is a legitimate clearance, but the remote-racer kind: its cost is that **custody is no longer individual**.

*Remark 28* (When the escape closes). The remedy holds only while the delegate is genuinely remote. If the delegate is a friend who happens to be visiting during the attack, *both* problems apply at once — custody was already non-individual, and the racer is now co-located and reachable — so the design becomes deadly-race-susceptible again. The user is thus forced into an awkward corner: trust someone they can rarely be near, and accept a perverse incentive to avoid physical proximity to a trusted friend.

## 5.2 A second time-lock design — BitVault

**BitVault** [19] is the same species. No authoritative specification is published: the vendor describes the product as “100% Open Source,” but its code link resolves to an organization with no public repositories (archived at <http://archive.today/2026.09.01-225418/https://github.com/BitVaulty>), and its documentation and audit are served through a gated document-sharing service. What we have is the vendor's pages, a concept whitepaper [20] that is self-attributed and hosted outside the vendor's own links, and correspondence of 26–27 August 2026 (all accessed 1 September 2026); they do not agree in every particular. This bounds what can be claimed and is *not* an allegation: we have no evidence that any material was withdrawn, and observed none at any earlier time.<sup>3</sup>

Taking that whitepaper as valid, the spending paths carry timelocks: a primary needing a custodian co-signature behind a configurable 2h–15d relative CSV, a one-year fallback on the user's own two keys, and a three-year custodian path, the last two with no gatekeeper. The policy gates the primary on `older()` and both fallbacks on `after()`, so the primary delay is a maturation from funding rather than a window on a request, and the multi-year locks do not re-arm. Both cases fail. Before the one-year date  $\mathcal{A}$  coerces  $A$  and  $B$  and kills, the only party who would tell the gatekeeper to withhold being  $\mathcal{V}$ : removal is what buys settlement, which is Lemma 1 exactly, and the pivotality is the victim's as *witness* rather than as racer, since the primary path needs the custodian's signature and  $\mathcal{V}$  cannot use it alone either. After the date he coerces  $A$  and  $B$  and spends at once, ungatekept — no delay, no notification, nothing to cancel.

Property 1 we leave open: the documentation assumes hardware and offers an optional “copy of  $A$  kept in a separate location,” but nothing in the construction forbids running it from memorized seeds and re-deriving on demand, which would be deviceless. Property 2 fails in both directions — the optional off-chain gatekeeper can deny a user the quick path by declining to release the custodian's co-signature, and past the three-year lock a custodian and that gatekeeper can spend together.

Property 3 we also leave open, for a reason worth stating generally: **obscurity does not advertise itself**, so classifying a design as obscurity-dependent is always an inference from what it would have to assume in order to make sense. Here the deadly-race exposure is plain on the published policy, plain enough that a Kerckhoffs-aware design would not ship it; the design therefore *suggests* an  $\mathcal{A}$  assumed ignorant of it — a reading the scarcity of public-facing documentation would fit. The only other reading we can construct is an  $\mathcal{A}$  assumed willing to coerce but not to kill — not a milder premise but a wager on his  $c$  (Definition 4), staked on

<sup>3</sup>The architectural description is the company's; the classification is the present author's, and is relative to this paper's threat model. In correspondence the company confirmed that the current configuration employs no autonomous, un-coercible watchtower and no pre-signed rescue sweep; that rejection of a pending authorisation requires intervention by the user, a separately positioned user-controlled device or signer, or a designated party; and that once the required signatures and timelocks are satisfied the resulting transaction cannot be reversed. It further agreed that the “publicly verifiable, making kidnapping attempts pointless” formulation was too categorical, and describes the decoy layer as plausible deniability rather than a guarantee.

the holder’s life. Undetermined is the honest entry between the two, and it costs nothing: the coercion result stands on structure either way.

A further caveat sharpens R2 itself, and a distinction sharpens the caveat. What matters is not only whether the delegate is remote, but what  $\mathcal{A}$  needs from them. Where the delegate’s required contribution is an *act* — a payment, a signature — coercion routes through  $\mathcal{V}$  to them and a humane delegate performs it. The act is self-verifying:  $\mathcal{A}$  knows when he has succeeded. The Balland case (§5.6) is exactly this failure, and it is the reality clause D8 idealizes away.

Where the required contribution is an *omission*, the analysis does not transfer. To win,  $\mathcal{A}$  needs the delegate to abstain for the length of the window, and abstention has no verifiable permanence: the delegate may comply visibly while retaining the means invisibly, and by NNPP (Assumption 4) cannot be made to demonstrate that no retained means exists — an undisclosed duplicate of a cancellation credential being, among the available evasions, close to the cheapest. The same negative existential that convicts  $\mathcal{V}$  in the individual case, that one cannot prove one has nothing left to give, acquits the delegate here, because it is now  $\mathcal{A}$  who needs the proof and cannot obtain it.

The incentive inverts with it. In the individual configuration, removing  $\mathcal{V}$  *wins* the race: the only party who could cancel is gone. Where the pivotal racer is unreached and may hold an undisclosed means of cancelling, removing  $\mathcal{V}$  *loses* it — the threat that keeps the delegate passive expires with the hostage, and the delegate then acts.  $\mathcal{A}$ ’s dominant strategy turns from removal to preservation, and he comes to need  $\mathcal{V}$  not merely alive but demonstrably so. This is the mechanical content of R2: not simply that the pivotal racer is out of reach, but that putting them out of reach makes  $\mathcal{V}$ ’s survival instrumental to  $\mathcal{A}$  rather than an obstacle to him.

The exchange should be recorded as an exchange. NNPP protects the funds by denying  $\mathcal{A}$  satisfaction, and it does so by relocating the pressure: an  $\mathcal{A}$  who cannot verify the delegate’s incapacity has reason to press the delegate. Shared custody does not dissolve the problem so much as move it to a party who is harder to reach and less exposed. That is a defensible trade, and it is why R2 is an honest route rather than a solved one.

### 5.3 The industry recommendation set

The two designs above are products; the same structure appears in the *standard advice*. Industry threat-intelligence reporting on wrench attacks states the objective in terms close to ours — custody should “remain resilient even when a human is under duress” — and recommends, among other controls, **withdrawal delays**, transaction-size thresholds, **staged vault architecture**, **emergency freeze** processes, and multisig or MPC with **geographically distributed signers** [6].

Criterion 1 does not reject this advice; it *partitions* it, and that partition is the framework’s practical content. Each of these controls interposes a delay or a veto, so each defines a competing racer, and the question Lemma 1 forces is only whether  $\mathcal{A}$  can reach whoever holds the release:

- Where the freeze or cancellation authority rests with a *remote, institutionally separated* party — a provider desk, an autonomous watchtower, a co-signer beyond reach — the control clears the bar by route R2, at that route’s standing cost: custody is no longer individual.
- Where the delay is one the *holder* can release, or the freeze authority is a device or a person on the scene, the control leaves  $\mathcal{A}$  a feasible-but-unfinished path and is therefore a race — the gray outcome the Criterion forbids, and the configuration in which the control is worse than its absence.

Geographic distribution of signers is the same statement made in space rather than in time (Definition 11), and inherits the same condition. The recommendations are thus sound *exactly when* the reachability question has been asked and answered — and it is that question, rather than the length of the delay, that separates a delay which buys rescue from a delay which buys  $\mathcal{A}$  a reason to eliminate the other racer.

#### 5.4 Self-destruct (wallet-erase) PINs, and the decoy beside them

The designs above interpose a delay or a veto. A large and widely shipped class does something categorically different: it *erases*. An alternate PIN, when entered, destroys the device’s copy of the secret. Blockstream’s Jade documents such a PIN under both of the names its help center uses for it — *duress* and *wallet-erase* — as an alternate PIN that “automatically delete[s] the encrypted wallet information from Jade if entered,” presented as “especially useful in the event of a physical threat, as you can provide this PIN to an attacker.” On entry the device reports only “Internal Error,” so that the erasure reads as a malfunction [21, 22].

Grant the design the strength it genuinely has, and which the delay architectures above lack. The erasure is *instantaneous and covert*:  $\mathcal{A}$  learns of it only once it is complete, so there is no interval in which he can intervene to prevent it. Nothing in what follows turns on  $\mathcal{A}$  out-racing the mechanism, because he cannot.

**The claim decomposes, and only one half is real.** What the design offers  $\mathcal{V}$  is not one mechanism but two, stacked and easily conflated:

- (a) *the erasure* — a genuine, irreversible state change, which happens whatever  $\mathcal{A}$  knows or believes; and
- (b) *the belief* that nothing further remains to be had — manufactured by the “Internal Error” report and the return to setup, which are engineered to read as breakage rather than as a move. This is a *presentation* choice and not a property of erasure as such; Remark 30 records a vendor asked directly for it who declined.

Only a is Kerckhoffs-independent. The whole question is how much coercion resistance it carries on its own, and the answer, from the vendor’s own instructions, is *none*.

The same documentation warns that afterwards “there is no way to regain access to your funds except by restoring using your recovery phrase,” and directs the holder to “make sure you have a backup before proceeding” [22]. The feature therefore *presupposes* a surviving backup: a holder who follows the documentation has one, and a holder who does not has destroyed the stash rather than defended it. What erasure removes is thus not the secret but the *device path* to it — and a path that cost  $\mathcal{A}$  nothing to lose, since  $\mathcal{A}$  leaves  $\mathcal{V}$  alive, present, and still holding everything needed to restore. In the ideal model  $\mathcal{A}$  contributes exactly nothing to the floor, for the same reason a stage below the memory-hard chain contributes nothing to it (Remark 22): what is reachable without work is not a barrier.

All of the design’s protective content therefore sits in b, which is an obscurity bet and nothing else. This is the configuration Criterion 1 forbids, and NNPP (Assumption 4) is what makes it hold:  $\mathcal{A}$  cannot verify that no further copy exists — and here he need not even reason his way to the suspicion, since the manual states that one should. He is therefore correct, by the design’s own construction, to read the erasure as the closing of one route rather than the end of the matter. The route that remains runs through  $\mathcal{V}$ ’s memory of a seed phrase, or of where a backup is kept: the erasure has deleted the path that required hurting no one and retained the path that does.

*Remark 29* (The obscure self-destruct is the obscure scheme with no winning branch). Obscure defenses are usually compared by *how likely*  $\mathcal{A}$  is to be ignorant. The sharper comparison is *how many* independent ignorances each one needs, and here the self-destruct is strictly the worse bet — worse, in particular, than the decoy it is commonly sold beside.

A decoy needs one premise: that  $\mathcal{A}$  does not know decoys exist. Grant him that ignorance and the decoy *terminates the encounter* — he leaves with a plausible balance and a story that closes.

A self-destruct needs two, and they are independent. It needs  $\mathcal{A}$  ignorant of the feature (sub-Kerckhoffs), so that he reads the failure as breakage rather than as a deliberate move; and it needs him ignorant of NNPP’s consequence (sub-NNPP), so that he does not reason to an undisclosed backup. The second is not implied by the first, and this is the design’s central weakness: *even total success at the first leaves the encounter open*. “This device is broken” does not entail “this

person has no bitcoin”; it entails that the instrument for reaching it is broken, whose natural sequel is the question *where is your backup?* — addressed to a  $\mathcal{V}$  who is still present and still knows.

So the self-destruct fails against an informed  $\mathcal{A}$  for the ordinary Kerckhoffs reason, and against an ignorant one it fails anyway, because the belief it manufactures does not discharge the encounter. It is the one obscure scheme with no branch on which it wins.

We did not have to derive that regress: a vendor documents it. Coldcard ships trick PINs whose actions include a duress wallet — “a personal safety feature. . . If you must reveal a PIN under duress, give the duress PIN instead of your Main PIN” — alongside *Brick Self* and *Wipe Seed* [23]. Its firmware documentation then addresses the case where the decoy is not believed, and its advice is to escalate to destruction: “you may be asked what the actual duress PIN is, while under duress. We suggest providing the ‘brickme’ PIN in that case” [24]. That is the manufacturer contemplating an  $\mathcal{A}$  who knows decoys exist, and recommending, as the answer, an irreversible erasure performed with  $\mathcal{V}$  still in the room and still able to be asked for a backup — the ignorance-count argument above, arrived at independently and published as guidance. The same document instructs the holder to “triple-check your backups are correct” before enabling that PIN, so the backup presupposition of §5.4 is not peculiar to one vendor either.

The same page goes further, in two admissions worth quoting because a vendor rarely makes them. The first concedes that the decoy is distinguishable to the  $\mathcal{A}$  who bothers: “if you are somehow facing an attacker who is willing to verify he has the real main PIN, it’s possible that careful analysis of system responses will imply he’s working with the duress PIN.” The remedy offered is not a repair but a request — “if you discover any sequences that reveal this easily, please tell us and we’ll see if we can cover them up better” [24]. That is obscurity stated as a maintenance programme, and it is the clearest statement we have found of what Property 3 rules out. What is being defended is not that the decoy is indistinguishable, but that no one has yet published how to distinguish it; the defence is administered by patching the disclosures as they arrive. A scheme whose security is the current state of the disclosure backlog is exactly the scheme whose guarantee  $\mathcal{V}$  cannot invoke in the room, because  $\mathcal{V}$  cannot know which sequences  $\mathcal{A}$  has read.

The second admission is the denial move itself, offered where bricking is unpalatable: “alternatively, you can say you set the duress PIN once, but have since forgotten it.” The escalation to destruction and this are the document’s two exits, and both are the naive tier. The claim *I have forgotten it* is a proposition about  $\mathcal{V}$ ’s own knowledge, and NNPP (Assumption 4) is precisely the observation that  $\mathcal{V}$  cannot close it: there is no move that converts a true *I cannot* into an  $\mathcal{A}$ -verifiable one, and by Definition 6 the scheme’s fate on that branch is decided by whether he believes it. The vendor has enumerated the branches and found, on each, a bluff.

That the line is published guidance sharpens rather than softens it, in the way Remark 33 describes for the coached decoy. A script printed in a manufacturer’s manual is a script  $\mathcal{A}$  can read as easily as  $\mathcal{V}$ ; the sentence that teaches the excuse also teaches that the excuse is taught. Its credibility declines in proportion to its circulation, and the holder who most needs it inherits the discount without ever having been told it exists.

*Remark 30* (The class, and what each vendor actually claims for it). Three vendors ship the pair, which is what makes this a design consensus rather than one firm’s misstep; but they do not make the same claims, and the difference is worth recording precisely. Blockstream markets the *erasure* for physical threat while claiming nothing of the kind for its passphrase wallets. Coldcard markets *both* halves, describing the duress wallet as a safety feature and the brick PIN as the fallback when it fails. Trezor is the mirror image of Blockstream: its passphrase wallets are sold on deniability, while its wipe code — which the vendor itself calls a “self-destruct” PIN, and which erases the recovery seed with it — is documented without any mention of duress, coercion or physical threat [25]. That vendor was moreover asked, in January 2022, to make the wipe *concealable*: a user-defined message on entry, so that an erasure could be presented as a hardware fault rather than announcing itself. The request was closed as *not planned*, and the device still

reports the wipe plainly [26]. Component  $b$  is therefore not intrinsic to the self-destruct class at all: it is a presentation one vendor builds and another, asked for it directly, refused. What §5.4 analyses under  $b$  is specific to the “Internal Error” design, and a wipe that announces itself carries  $a$  alone — which, by the argument above, is no coercion resistance at all, and is not sold as any. Two of the three therefore decline to make the coercion claim on one of their two mechanisms. We note it because the criterion is not advanced against vendors, and a vendor that ships a mechanism without claiming it defeats coercion has not made the error this section is about.

*Remark 31* (Two bluffs the erasure invites). An  $\mathcal{A}$  who knows the design (Assumption 2) need not walk into the erasure at all. Two alternatives stand open to him, each a wager rather than a certainty:

1. *Take the device hostage up front.* Rather than enter any PIN, seize the device and hold it as a hostage token (§3.7, Remark 17). An erase PIN is irrelevant to a coercer who never authenticates; and by Remark 18 the terms on which  $\mathcal{V}$  ransoms the device are themselves a probe for the very backups  $\mathcal{V}$  cannot disprove.
2. *Coerce against deployment.* Brutalize  $\mathcal{V}$  to deter the PIN’s use in the first place, wagering that a holder under sufficient duress will not spend the one irreversible move available to them.

Neither wager is sure to pay, and that is precisely the point: both are bluff games, and the holder who has bet on an erasure has entered one without holding the means to end it. A  $\mathcal{V}$  who erases and then *truthfully* denies retaining any backup arrives at the position the registry already records — the streamer whose denial was simply disbelieved, and who was tortured for hours regardless [5] (§5.6).

The “Internal Error” disguise inherits the same objection from a different direction: it is a bet on  $\mathcal{A}$ ’s ignorance of a consumer product, and it degrades exactly as that product’s own public documentation becomes known (§3.6). The class therefore lands in the **naive** tier of Definition 9 precisely as that definition describes it — an obscure self-destruct, safe only against a sub-Kerckhoffs attacker — and it is worth keeping the self-destruct distinct from the *cancel* key named alongside it there. A cancel reverses a pending action and thereby names a second racer; a self-destruct destroys a path and names none. Only the former is subject to Lemma 1, which is why the two fail differently despite sharing a tier. Remark 16 has already disposed of the tier as a whole, by the dilemma that an  $\mathcal{A}$  whose ignorance can be assumed makes the elaborate mechanism unnecessary, and one whose ignorance cannot makes it ineffective; the present subsection is that dilemma instantiated in a shipping product.

**The companion decoy, classified separately.** The same device ships a second and quite different obscure mechanism, and the two should not be run together. A BIP-39 passphrase “is added to your Jade’s recovery phrase or SeedQR and is used to create your wallet”; “any number of passphrases can be entered to change the wallet you will log into”; and the device “will not store your passphrase” [27]. Distinct passphrases therefore yield distinct wallets from one recovery phrase, and the familiar community practice is to keep a low-balance wallet at the empty passphrase to be surrendered under coercion.

Two observations. First, and to the vendor’s credit, *Blockstream does not market it this way*: the passphrase documentation claims no deniability, decoy or duress use, in pointed contrast to the erase-PIN pages, which name physical threat explicitly. The duress framing here is the community’s, not the manufacturer’s, and the distinction is worth preserving in any account of who claimed what.

Second, the mechanism is the plausible-deniability layer already treated in §3.6 and dismissed for the same Kerckhoffs reason: it rests on  $\mathcal{A}$ ’s ignorance, and an  $\mathcal{A}$  who knows that passphrase wallets exist simply persists past the surrendered one, with NNPP guaranteeing  $\mathcal{V}$  cannot close the question. It is not, however, the *same* failure as the self-destruct’s, and by Remark 29 it is the less bad of the two: the decoy at least has a branch on which it wins, since an  $\mathcal{A}$  who believes it has been paid and can leave. The self-destruct has none. That a single product ships both,

marketing only the weaker one for physical threat, is the clearest illustration in this section of how little the design space has been mapped against the Criterion.

*Remark 32* (A vendor-documented instance of the naive tier). The decoy’s duress framing was the community’s, not the manufacturer’s. One vendor does make the claim in its own documentation. Sparrow’s FAQ, answering “How can I choose where my wallet is saved?”, documents a data-directory flag and adds that “[t]his feature allows you to store all Sparrow data on removable media making for more plausible deniability” [28].

The feature does not guard the asset. Sparrow’s own best-practice page has the seed words backed up separately from the device, “ideally in a different location” [29] — so the concealed directory is neither necessary nor sufficient for  $\mathcal{K}$ . A Kerckhoffs  $\mathcal{A}$  coerces the seed backup, loads it into a wallet of his own, and never touches the concealed medium at all. What is left for the medium to do is carry the holder’s belief that, having surrendered what was found, they may now say there is nothing more. The decoy at least stands between  $\mathcal{A}$  and the coins; this stands nowhere near them.

That is Definition 6 exactly — a guarantee contingent on  $\mathcal{A}$ ’s ignorance, here of a *documented* feature — and the continuation is the one the registry records. A denial offered into the prevailing norm is pre-discredited (Corollary 1), so the rational continuation of an interrogation that opens with one is more coercion, and the truthful denier is brutalized alongside the lying one (§5.6). The feature adds nothing but the confidence to make the denial.

*Remark 33* (The decoy is sold, and coached for credibility). The decoy is not merely a practice the community tolerates; it is a product and a curriculum. A Spanish-language self-custody consultancy lists, among paid service tiers, “*Passphrase (tu bóveda señuelo)* — Billetera señuelo + billetera real. Protección ante coacción/coerción” — a decoy wallet marketed, in as many words, as protection against coercion [30]. Published guidance goes further and coaches the performance: one guide has the holder keep a small balance on the bare seed as “what you’d hand over under physical coercion”, and adds that “for this to work, the decoy balance must be *believable* — enough to look like a real wallet, not enough to devastate you if lost” [31]. The counsel is also delivered live: one of us appeared as guest on a self-custody programme [32] in which the host worked through a decoy-construction walkthrough over the guest’s contrary advice.

That last instruction is the naive tier stated as a training objective. A decoy need be believable only to an  $\mathcal{A}$  who might disbelieve it, and the counsel is silent on what happens when he does — which, by Definition 6, is the whole question. A better-rehearsed lie raises the chance  $\mathcal{A}$  takes the bait at the margin; it does not create a branch on which the scheme wins against an  $\mathcal{A}$  who knows decoys exist, and NNPP (Assumption 4) guarantees  $\mathcal{V}$  cannot close that question however well the performance goes. What is sold as a mechanism is a wager on the attacker’s ignorance, priced and taught as though it were a control.

The counsel is also self-defeating in aggregate. Publicised training raises the *prior* that a fluent denial was rehearsed: once  $\mathcal{A}$  knows credibility is taught and sold, fluency stops being evidence of truth and becomes evidence of coaching. The remedy degrades the thing it sells — and not only for its customers. A holder who never trained now denies into an  $\mathcal{A}$  who discounts fluent denial generally, so the cost falls on people who bought nothing. This is the Obscurity Paradox (Remark 34) run backwards: there, a public non-obscure design makes  $\mathcal{V}$ ’s claims *more* plausible; here, a public obscure training makes them *less*.

*Remark 34* (The Obscurity Paradox — non-obscure designs strengthen obscure ones). TGPO enters neither bluff. There is nothing to erase, the deep secret having never been on the device (Theorem 3), and so nothing whose non-deployment can be coerced. But the comparison yields more than the absence of a failure mode, and the surplus runs in a direction the obscurity argument would not predict.

An obscure defense asks  $\mathcal{A}$  to believe a *claim* — that this is the only wallet, that nothing further remains — against both an incentive to disbelieve it and, by NNPP, no means of settling

it. A non-obscure defense changes what making that claim costs, and it does so by the mechanism set out in Remark 23: an  $\mathcal{A}$  who credits that an effective mechanism is in place is crediting that lying is not  $\mathcal{V}$ 's only available tool, and a holder with effective alternatives has correspondingly less reason to resort to one.

The ordering this inverts is worth stating plainly, and we name it **the Obscurity Paradox**: *a publicly specified, Kerckhoffs-clean design improves the standing of precisely those obscure denials it was built to make unnecessary*. Obscurity remains unsound as a primary defense, for the reasons §3.6 gives. It is not, however, indifferent which design it accompanies — a denial made from within a mechanism that removes the need to lie is a stronger denial than the same words spoken from within one that requires it. The same asymmetry has an epistemic cost, taken up next (Remark 37).

*Remark 35* (The Denial Spiral). The section has now derived two distinct ways a custody practice can worsen its holder's position, and the first deserves a name as the second has one.

Denial credibility is a *common resource*: it is supplied by the population of holders and consumed by each holder who denies. Under the adversary of Assumption 2, a norm of concealment and coached denial depletes it. Where denial is the expected script a denial carries no information, so an interrogation that opens with one yields  $\mathcal{A}$  no exculpatory update and its rational continuation is further coercion rather than release (Corollary 1). The individual holder's rational response to that environment — conceal harder, rehearse better — depletes the resource further, for everyone, including the truthful denier who bought nothing (Remark 33).

We call this the **Denial Spiral**. It is the Obscurity Paradox (Remark 34) run in reverse, and it is the companion of the Deadly Race rather than a special case of it: the Race prices the *terminal* move, the Spiral prices the *duration* — the abductions, long holds and torture the registry codes (§5.6). A design can be innocent of one and guilty of the other, which is why they are named apart.

**Corollary 5** (The Spiral is not a sixth axis either). *Whether a design depletes the denial commons is equivalent to its failing Property 3. A design whose operation requires the holder to deny — a decoy, a concealed erasure — adds a scripted denial to the population and lowers what any denial is worth; a design that does not, and whose specification is public, does the reverse, since a denial made from within it is consistent with an ordinary setup rather than being the script (Remark 34). Depletion therefore tracks non-obscurity exactly, and a Denial Spiral column would reproduce that one as surely as a no-hostage column would reproduce Property 1.*

*Exposure is a different matter, and it is not a property of any design. The Spiral's cost is ambient: it falls on whoever must be believed, including the holder of a design that depletes nothing, and including the person who, on his own account, held no bitcoin at all — the RATIRL case (§5.6), and the reason the externality is worth naming. The qualifier is owed: §5.6 records that the denial may itself have been defensive, and the externality argument is the branch on which it was true. No column can carry that, because no column has a row for the people it lands on.*

*Remark 36* (What kind of claim these are — a praxeological reading). Neither the Race nor the Spiral is an empirical hypothesis awaiting corroboration, and it is worth saying plainly why, because the alternative reading invites a demand this domain cannot satisfy.

Both are deduced from the fact of purposive action under a stated adversary: *given*  $\text{WDY}[W, c]$ , these are the conditions under which coercion is prolonged or the holder is removed. That is the posture Mises named *praxeology* — theorems derived from action rather than induced from observation, with history able to illustrate a theorem but never to verify or falsify it [33]. We claim the posture, not the whole method. The decision rule here is quantitative, and a strict praxeologist would resist it:  $p_A$  is what Mises called a *case* probability — the likelihood of this attacker prevailing in this encounter — and he denied that numbers belong there. The concession costs little, because nothing in the argument requires  $p_A$  to be *computable*; only that the ordering it expresses exists. What carries over is the epistemology, and that is the part that governs how the results should be read.

No observation could settle either claim, and the reason is internal to the subject. The only witnesses to an enacted Deadly Race are an attacker with every reason to deny the motive and a holder — or, where the Lemma’s prediction is realised, a family — with every reason to deny the asset. Asking a convicted murderer whether he killed to win a race for a stash whose existence the survivors also deny is not a weak instrument but a void one: both answers are the negative that NNPP (Assumption 4) holds cannot be disproved.

The Spiral fares no better under the same test. Ask a convicted kidnapper whether the coercion he prolonged did in fact reach a stash whose existence his victim still denies: an affirmative confesses a robbery on top of the abduction, and a *truthful* denial is indistinguishable from the answer a successful robber would give — while the victim’s own denial, which is what would corroborate either, is once more the negative that cannot be disproved. §5.6 works that fork through on the one case where it can be posed concretely: three readings stay simultaneously plausible, and each indicts obscurity. Neither witness can be believed, and no third party was present. The instrument that would test the model is disabled by the very mechanism the model describes.

The standard that does apply is *conditional soundness*, and it is checkable without data. The theory can be shown wrong, in two specific ways. **Break the derivation:** exhibit a reachable state that is sufficient for  $\mathcal{K}$  but unfinished, with no access-dominating delegate, in which  $\pi = 0$ . **Break the *a fortiori* step:** show that the modelled adversary is, in some respect that bears on the result, *weaker* than a real coercer — for a bar cleared against an adversary granted more than reality gives him (§3.1) is cleared against the real one only while that grant holds. Neither route requires a case, a victim or a survey; both are settled by reading the argument.

What is genuinely undecidable is only the per-case attribution — whether *this* killing was enacted for *that* payoff — and the paper never asserts it. This is also why the output is a **bar** rather than a finding. Whether a protocol leaves a sufficient-but-unfinished state is decided by reading the protocol; whether a norm depletes denial credibility is decided by reading the counsel. Neither requires a victim, and that is the point of the whole construction.

*Remark 37* (Obscurity anosognosia — why the crisis resists diagnosis). Four censoring mechanisms compound. Three remove the evidence that would settle whether the Lemma’s prediction is being enacted in the field; the fourth removes the response to the evidence, and so survives its publication.

The first is the one this section has already named: obscurity does not advertise itself. No vendor documents a scheme as depending on  $\mathcal{A}$ ’s ignorance, so for any given victim the question *was this holder running an obscurity-dependent design?* has no recorded answer. The exposure variable is not merely unmeasured but unmeasurable from public sources, which forecloses the obvious epidemiological question — what do the victims of prolonged coercion, or of the terminal move, have in common?

The second is recorded in §5.6: a homicide is reported as a homicide, not as a robbery whose victim was then killed, so the enactments most diagnostic of the Lemma are the least likely to enter the registry at all.

The third forecloses the one route by which the first two might have been worked around. Where  $\mathcal{V}$  is killed, the people who survive them are the only remaining source for what that holder was running — and they are, by construction, the party least able to supply it. A scheme that depended on  $\mathcal{A}$ ’s ignorance depended on  $\mathcal{V}$ ’s discretion, and discretion is not selective: the family was not told either. Where they do know, disclosure carries no benefit to them and an evident cost — it advertises that a household which held bitcoin may hold it still — so it is withheld unless it might help resolve the case. The one person who could have answered the question is the one Lemma 1 predicts will not survive to be asked.

The fourth is of a different kind, and is why publishing the first three does not dissolve them. Conceding the argument is expensive to whoever concedes. A vendor who grants that a shipped feature is worse than its own absence has said something about products already in customers’ hands; a trainer who grants that coached denial depletes the credibility of all denials has said

it about a service already sold; a holder who grants it has said it about a setup they are still running. The cost of conceding is private, immediate and borne by the conceder; the cost of *not* conceding is diffuse, deferred and borne by holders who are not in the room. That is the same commons structure as the Denial Spiral (Remark 35), one level up: there the depleted resource is the credibility of denials, here it is the field’s capacity to correct itself, and in both the individually rational move is to wait for someone else to go first. Recalcitrance is cheapest exactly where it is costliest in aggregate.

The precedent usually reached for is Semmelweis, and it repays being stated rather than gestured at. In 1847, after his colleague Jakob Kolletschka died of a wound taken from a scalpel during an autopsy and showed the pathology of childbed fever, Semmelweis inferred that the fever was carried on the hands of physicians arriving from the dissecting room, and imposed washing in chlorinated lime; mortality in his clinic fell from roughly 18% to roughly 2% within the year [34]. The inference was available two decades before the germ theory that would explain it, and it was rejected anyway — not on evidential grounds alone but on positional ones, since accepting it meant accepting that physicians’ own hands had carried the disease [35]. That is the narrow aptness, and it is the same shape as the mechanism above.

The wider instruction cuts against us, and we take it. When the 1861 monograph failed to persuade, Semmelweis turned to open letters denouncing his critics as murderers. The denunciations did not shorten the interval before adoption; Pasteur’s germ theory and Lister’s antisepsis did, two decades on, and the interval was measured in lives [35]. It is one reason this paper states a bar rather than an indictment, classifies designs rather than people, and credits vendors where the record credits them (§5.4).

That restraint is also what keeps this remark from becoming unfalsifiable. A mechanism that predicts silence will be read as agreement is a licence, not a mechanism, and we decline it: the claim is that concession is *costly*, not that any particular non-response is *concealed assent*. Silence has ordinary causes — time, authority, counsel, disagreement never written down — and nothing here licenses an inference about any individual who did not reply. What it predicts is aggregate: that the availability of the analysis will not by itself produce its adoption, and that adoption arrives first wherever conceding is made cheap — a neutral weakness class to cite, a reviewed venue to cite it from, a bar defined before anyone is measured against it.

So the covariate is censored by what vendors do not say, the numerator by how the terminal event is classified, what remains by what the bereaved do not know or will not repeat, and the correction by what it costs the corrector. The structure is that of *anosognosia* — the dementia patient whose disease erodes memory and, with it, the capacity to notice that memory is eroding. One cause, two effects, and the second conceals the first. Obscurity is simultaneously what makes a design fail under Kerckhoffs and what erases the record by which the failure would be recognised. We draw no empirical claim from this: we cannot, which is the point. What it establishes is that the diagnosis is not available from the data by waiting, and that a field which waits for it will keep not finding what it is not positioned to see. That is an argument for the formal route — the bar has to be derived, because it will not be measured.

## 5.5 TGPO, and the honest comparison

TGPO clears by the *other* route: the secret is tacit and never on the device (Assumption 3, Theorem 3), so a seizure yields nothing feasible — there is no racer to reach, and the time-lock is only a delay *on top of* an already-clean failure. Read through the Four Properties (Corollary 3), the comparison is a covering-set statement: each shipping design buys (or tries to buy) coercion resistance by *trading at least one* property TGPO keeps — and one of them trades a property and *still* fails. TGPO is the only design holding all four.

| Design            | (1) Know. | (2) Indiv. | (3) Non-ob. | (4) Coerc. | DR-safe | What it trades                                                                           |
|-------------------|-----------|------------|-------------|------------|---------|------------------------------------------------------------------------------------------|
| TGPO (this work)  | ✓         | ✓          | ✓           | ✓          | ✓       | holds all four (tacit; R1)                                                               |
| Rewind Bitcoin    | ✗         | ✗          | ✓           | ✓          | ✓       | individual custody: a remote delegate (R2)                                               |
| BitVault          | ?         | ✗          | ?           | ✗          | ✗       | individual custody both ways — denial blocks the user, collusion spends past three years |
| Casa              | ✓         | ✗          | ✓           | ✓          | ✓       | custody shared with remote co-signer                                                     |
| Self-destruct PIN | ✗         | ✓          | ✗           | ✗          | ✓       | non-obscurity: the erasure must read as malfunction (Coldcard, Jade)                     |

Column (1) is the covering point: TGPO is the *only* knowledge-based (deviceless — the secret is knowledge, never a device or a location) row, because every competitor keeps a *transmissible* secret on a device. The fifth property of §3.7 adds no column: by Corollary 4 it is *equivalent* to that one, so it would reproduce it cell for cell. Neither does the Denial Spiral, for the reason Corollary 5 gives.

The **DR-safe** column is not derivable from the other four, and that is why it is there. It asks the criterion directly — *can* a single encounter leave a state of positive pivotality? — and the answer does not track the properties: **Rewind** forfeits two of them and is DR-safe nonetheless, clearing by R2, while **BitVault** forfeits them and is not. The **self-destruct PIN** is the instructive row: it is DR-safe and still condemned, because a self-destruct destroys a path and names no second racer, so Lemma 1 does not reach it — its failure is non-obscurity, a different one. The column is worth reading for what it does *not* say: four of the five shipping designs clear the bar this paper defines. The claim was never that every deployed countermeasure is lethal. TGPO’s mark is subject to the bounded residual conceded in §4.7. **Rewind** trades individual custody for a remote delegate (R2). **Casa**’s collaborative multisig — the client’s keys on separate hardware plus a Casa co-signing key — trades individual custody: the co-signer is a remote party, and a coercion-resistant delegate is remote *by definition*, so the cost is individual custody, not non-geography (Casa asks the user to visit no site of its own) [36]. **BitVault** trades individual custody for the same reason Casa does — a required co-signature on its primary path — but without obtaining what the surrender is meant to buy: it is the one row that pays a property and, per §5 above, still lands in the gray zone. Its remaining two columns are left undetermined, for the reasons given there. The axis tilts further to R1 for a blunter reason (Remark 9): the delegated rows are not self-custody at all — their clearance assumes a delegate that refuses even the legitimate owner — and, the rescue capability being *transmissible*, they keep a clean-loss exposure the tacit route does not. *Not your keys, not your coins.*

## 5.6 Empirical grounding in the registry

The two stipulations that lean on the physical-attack registry [2] — the coercion-time budget  $t$  (§8) and the murder-cost weighing of the residual-window analysis (§4.7) — admit directional empirical support. The registry records each incident only as *date*, *victim*, *location*, *headline*; it carries no duration or outcome field, so we code the 352 incident headlines (at commit 9a4a62a) by keyword. The coding is coarse and the categories overlap; a reproducible script accompanies the paper.<sup>4</sup> Several things stand out.

<sup>4</sup>registry-analysis/ in the source tree gives the exact keyword patterns, per-year counts, and caveats.

*The attack is frequent and worsening.* Recorded incidents climb from a handful a year before 2017 to a peak of 85 in 2025 — the growth the introduction cites, now counted rather than asserted.

*The terminal move is real.* At least 5.1% (18/352) of headlines record a *killed* or *murdered* victim. A homicide is usually reported as a homicide, not as a “Bitcoin attack”, so this is a *lower* bound — yet even so it shows the assassination outcome the Deadly-Race Lemma (Lemma 1) predicts is attested in practice, and that the cost  $c$  is weighed against an action attackers demonstrably take.

That censoring compounds with a second: obscurity does not advertise itself, so no victim’s exposure to an obscurity-dependent design is on record either. The numerator is censored by how the terminal event is classified and the covariate by what vendors do not say, which is why the pattern cannot be read off the data and the bar has to be derived instead (Remark 37).

*A separately sourced count agrees on the trend.* The registry is a media-sourced convenience sample, so an independent base matters. Industry threat-intelligence reporting — assembled from law-enforcement disclosures, court documents, interior-ministry statements and victim testimony rather than from the registry — records 72 verified incidents in 2025, up 75% year on year, with confirmed losses of \$40.9M that the report itself calls an undercount [37]; and 52 incidents in the first half of 2026 against 39 a year earlier, with losses of \$124.1M against \$10.5M [6]. Stricter verification yields counts below the registry’s headline sweep, as expected; the direction, the European concentration, and the category mix agree. Two features bear directly on this paper. Reported *home invasions* rose from 1 to 20 of 52 incidents year on year — coercion is moving to where the holder lives, which is the setting the model assumes. And their *confirmed* homicide count is far below our headline coding (1 in the half-year, against our 18 over twelve years), which is what two different evidentiary bars should produce and which reinforces that the share coded here is a floor, not an estimate.

*Short attacks dominate; the registry over-counts the dramatic.* Most wrench attacks are *hand-over-now* events — a threat, an immediate transfer, minutes on scene; the long, cinematic kidnapping is the exception. A media-sourced registry inverts this, because a mutilation-kidnapping is news where a knife-point mugging is not, so the kidnap share coded below *over*-states the long tail rather than establishing it. That TKBA is realistically *less* effective against a long-duration attack — where hours of pen-and-paper effort under coercion *might* walk a deep stage — is granted, but it is *speculative* pending the field study of §8. What is *not* speculative is that the fielded alternatives fail; the cases below make that concrete.

| Signal (headline keyword coding)                  | Incidents | Share |
|---------------------------------------------------|-----------|-------|
| Killed / murdered victim ( <i>a lower bound</i> ) | 18        | 5.1%  |
| Kidnap / hostage / ransom (long-hold)             | 139       | 39.5% |
| Armed robbery at gunpoint (short-hold)            | 101       | 28.7% |
| Home invasion / break-in / raid                   | 40        | 11.4% |
| Torture or physical violence                      | 82        | 23.3% |

Categories overlap (a kidnapping may also be violent or fatal) and do not sum to 100%; shares are of all 352 incidents, and the modality labels are a headline proxy, not a measured duration. The sample is media-reported and survivorship-biased — deterred, unreported, and silently-fatal cases are absent — so every figure is directional and the fatality share in particular is a floor. Remark 37 gives the two mechanisms that censor it, and why they leave the diagnosis unavailable rather than merely imprecise.

**The mix is not universal, and the Lemma’s bite tracks it.** The global shares above average over jurisdictions whose compositions differ by an order of magnitude. Taking the ratio of long-hold to short-hold codings,  $K:G$ :

| Jurisdiction              | $n$ | Kidnap | Gunpoint | $K:G$ |
|---------------------------|-----|--------|----------|-------|
| Brazil ( <i>small n</i> ) | 12  | 75.0%  | 8.3%     | 9.0   |
| France                    | 61  | 57.4%  | 8.2%     | 7.0   |
| All incidents             | 352 | 39.5%  | 28.7%    | 1.4   |
| United States             | 59  | 28.8%  | 40.7%    | 0.7   |
| Russia ( <i>small n</i> ) | 10  | 30.0%  | 50.0%    | 0.6   |

The comparison that carries weight is France against the United States: the two largest national subsets, of near-identical size (61 and 59), with *inverted* compositions — a tenfold difference in  $K:G$ . This is the sharpest available check on the media-inflation caveat just stated. That caveat explains a *level* — why the registry over-counts dramatic incidents everywhere — but it cannot explain this *ordering* without the further claim that French reporting is ten times more kidnap-selective than American, in a market that is at least as heavily covered. The parsimonious reading is that modality composition genuinely varies with local conditions — firearm availability, the presence of organized kidnap-for-ransom capability — as it does for kidnapping outside this domain entirely.

Two confounds remain and are not dismissed. Outside France and the United States the subsets are too small to carry an argument, so the other rows are directional only. And the coding is applied to English-language headlines, several of them translations, so a systematic vocabulary difference — French sources reaching for “kidnapped” where American ones reach for “robbed at gunpoint” — would produce part of this spread without any difference in the underlying events.

If the composition is real, it matters for scope rather than for validity. The Deadly-Race Lemma (Lemma 1) is a statement about the *released* holder, and release is the trigger (Remark 7); a hand-over-now robbery at gunpoint approximates the clean  $p_A = 1$  endpoint — the attacker takes what is reachable and leaves — whereas a multi-hour hold ending in a deliberate release is exactly the gray outcome. Coercion risk is therefore not merely higher in some jurisdictions but *differently shaped*, and the No-Gray-Area Criterion (Criterion 1) does its work where long holds dominate. That is the regime the record is currently growing into: France alone accounts for a majority of the incidents recorded in 2026.

**The attacker often does not need the holder at all.** A fourth signal falls outside the keyword table because it was coded by hand: in 12 of the 352 incidents (3.4%) the person seized was *not* a holder but a relative held to compel one — a parent, a spouse, an adult child.<sup>5</sup> Two features of this subset matter more than its size.

It is *recent and concentrating*: eleven of the twelve fall in 2023 or later, and *nine* occur in a single jurisdiction (France) whose recorded incidents rose from about one a year through 2024 to 22 in 2025 and 34 by August 2026. Whatever share of that rise is genuine incidence rather than reporting intensity, the tactic is not a historical scatter; it is being adopted.

It is *not* the hostage token of §3.7, and the difference matters. There the leverage is a *seizable object* the protocol requires, which devicelessness removes by leaving no such object; here it is a *person*, whom no custody design renders unseizable. What this subset shows in the field is clause D8 failing: such an attacker needs no cryptographic access and no cooperation a protocol could withhold — only leverage, and a relative supplies it. This bears directly on the delegated route (R2), whose clearance rests on the deciding party lying beyond  $\mathcal{A}$ ’s coercive reach. The delegate a holder is most likely to actually appoint — a spouse, a parent, an adult child — is precisely the party this subset shows being seized. The route is not thereby refuted: it is bounded. A delegate who is reachable is not a second racer but a second victim, so R2 clears the bar only

<sup>5</sup>Keyword screening surfaced 16 candidates; each was read individually and four were discarded as false positives, including one in which “son of” named a *perpetrator* rather than a victim. **registry-analysis/** carries the screen and the full audited list. Because the coding reads a one-line headline, an incident that does not name a relative is no evidence that none was involved: 3.4% is a floor, and the coding can only understate it.

for delegates whose unreachability is structural rather than assumed — and that requirement, not the arrangement’s convenience, is what the data prices.

**A case in point — Balland.** The January 2025 abduction of David Balland, a co-founder of the hardware-wallet maker Ledger [38], is instructive precisely because it is an *embarrassment to the self-custody creed*. A professional champion of self-custody was seized from his home and mutilated — a finger severed to force payment — and neither wealth nor prominence prevented it. He was freed by an elite *state* police unit, and the ransom — paid not by Balland but by an associate — was clawed back largely because part of it moved in a *centralized* stablecoin whose issuer could freeze it. A self-custody pioneer made whole by the state and by centralization: two mechanisms external to, and in tension with, the creed. Only the issuer freeze is independently verifiable; the remainder rests on the parties’ accounts, according to which roughly 5% was never recovered. The case also indicts the delegated route (R2): an associate *did* pay, which is exactly the compliance clause D8 assumes away — a real instance of the indirect coercion, routed through the victim’s body to a third party who will not refuse. For a brief hand-over-now robbery a remote delegate is unreachable in time and so seems to help; but merely *not carrying* the device — leaving the wallet in a home vault — does the same, while for the long attack the delegate is not out of reach but is the ransom channel itself. The lesson is not that Balland erred: it is that, against a determined long-hold attacker, the fielded options collapse to state rescue and centralized freezing, whereas the tacit route removes the leverage — no standing device, no hostage token, no delegate to pay — rather than relying on ex-post recovery.

**A case in point — RATIRL.** The registry coded above records the obscurity route’s failure directly. In one of its 2026 incidents (Sweden, 13 January) [5] a League of Legends streamer was bound and tortured for hours by armed intruders demanding cryptocurrency he denied holding; unable to credit the denial, they persisted long past any value the robbery could rationally return, leaving finally with only his Pokémon-card collection. This is Corollary 1 in the field: where the community norm is “deny that you own Bitcoin,” a coerced denial is *pre-discredited* — presumed to be the denial-script — so a Kerckhoffs-educated  $\mathcal{A}$  gets no exculpatory update, and the rational continuation of an interrogation that opens with a denial is *more* coercion, not less. The case fixes the *sign*: obscurity does not merely fail to protect — it can convert a robbery into a torture — *negative* coercion resistance, not a mere absence. And it does so on *every* reading of the denial, which forks three ways — each indicting obscurity. (i) If the denial was a defensive lie (there was a stash), secrecy still did not avert the assault — a textbook Kerckhoffs failure — and his only residual recourse is to *double down* on it and wager against a repeat. (ii) If it was true (no stash), he bore obscurity’s *negative externality*: a genuine non-holder attacked only because real holders are coached to look exactly like him — the very coaching that collapses the statistical holder/non-holder distinction. (iii) If the truth is in between — a real but smaller-than-presumed stash, too small to be worth the judicial exposure the attack incurred — *both* pathologies operate at once: others’ obscurity redirected the attack onto him, and his own failed to stop it. That all three are simultaneously plausible is itself the verdict: in the aggregate, across a growing population of holders and lookalikes, each is already occurring — and the obscurity culture, exactly as Kerckhoffs predicts, prevents none of them.

## 6 Quantum note

---

Memory- and depth-hardness is a tunable lever, and even its modest settings suffice here. A quantum search over an  $n$ -bit space needs  $\sim 2^{n/2}$  *coherent* evaluations of the oracle [39], and each such evaluation is a deeply sequential Argon2d computation whose full *depth* the coherence must survive — a cost the quantum resource-estimation literature already finds decisive even against oracles far cheaper than ours [40]. Even a *few-minutes*  $H^*$  per element therefore places Grover out of reach; an *hours-scale*  $H^*$  does so *a fortiori*. Hence sub-256-bit standards (96 or 128 bits per

stage) are quantum-defensible *without* a 256-bit seed — the memory- and depth-hardness behind the floor’s *hard* bound ( $64^*$ ) classically also caps the quantum speed-up.

*Remark 38* (The one bound Grover does dent: the live- $o_N$  residual). The duality’s *cheap* bound is a different matter. The floor also clears  $64^*$  via a cheap  $2^{96}$  ledger brute of  $Sh_N$  (Stipulation 1, S3), and that cheap side inherits *no* depth-hardness to guard it: Grover halves it outright to  $\sim 2^{48}$ . Concretely, a Quantum WDY attacker landing on an imminently- or currently-live final board  $o_N$  — the same coincidence that opens the tolerated residual window (§4.6) — faces only the cheap  $2^{96}$  brute of the terminal secret, and Grover reduces that to a  $\sim 2^{48}$  (“48-qubit”) floor.<sup>6</sup> This touches nothing but the residual: it does *not* reach the memory-hard  $64^*$  bound (guarded above), and it applies only inside the already- $\varepsilon$ -bounded live- $o_N$  sliver, whose incentive-safety argument (§4.7) turns on the kill economics, not on the brute’s width — a  $2^{48}$  Grover search buys  $\mathcal{A}$  nothing a classical  $2^{32}$  brute in that same window does not already grant.

## 7 Related work

**Coercion resistance.** The term is due to Juels, Catalano, and Jakobsson [41], in electronic voting: a scheme is coercion-resistant if a voter *cannot prove* to a coercer how they voted, so the coercer cannot verify compliance and the reason to coerce collapses. Our sense is different and complementary. There the protected act is a *choice* the coercer wants to dictate and safety comes from *unprovability*; here the protected object is a *bearer secret* the coercer wants to extract, and the danger is not that compliance is unverifiable but that a released holder remains a racer worth killing (Lemma 1). We inherit the word, not the model.

**Secrets recognized but not articulated.** The tacit premise (Assumption 3) — a secret its holder can *recognize* but cannot *transmit* — has direct antecedents. Bojinov, Sanchez, Reber, Boneh, and Lincoln [42] train an authentication secret by *implicit learning* so that even a cooperating, coerced user cannot articulate it (“rubber-hose” resistance); this is the closest prior instantiation of our TKBA, and TGPO can be read as carrying the same idea from an authentication token to a self-custody key behind a memory-hard time-gate. The broader phenomenon — identity carried in a behavior that resists articulation — is old: keystroke dynamics treat idiosyncratic typing rhythm as a biometric [43], and WWII traffic analysis identified individual wireless operators by the “fist” of their Morse keying long before either was formalized [44].

**Deniability versus unprovability.** Deniable encryption [45] and the duress-PIN / panic-password schemes it inspired let a coerced user *plausibly lie* — surrender a fake key, or open a decoy. NNPP (Assumption 4) is the sharp negative counterpart: rather than proving a false positive (“this is the key”), the holder would need to prove a true negative (“no other key exists”), which cannot be done. Deniability offers a lie the coercer cannot disprove; we show the holder cannot furnish the assurance the coercer would need, which is exactly why the coercer’s only remaining lever is the victim (§3). Note too (§3.6) that decoy and duress schemes rest on the adversary’s ignorance and so fail Kerckhoffs.

**Time-gates.** The construction’s wall-clock gate is a time-lock puzzle in the sense of Rivest, Shamir, and Wagner [16]; verifiable delay functions [46] are the modern, publicly checkable form of the same primitive, and either may instantiate the outsourced time-lock the orbit optionally uses (§4).

**Kidnap economics and marginal deterrence.** The Lemma’s shape is not new, and it would be wrong to present it as though it were. Selten’s kidnapping game [47] is the canonical formalization of the encounter we model, and the criminological literature on *marginal deterrence* — Stigler [48], formalized by Shavell [49] — has long carried the observation that an offender facing

<sup>6</sup>Each Grover attempt is one cheap evaluation of the terminal step — a single  $H(o_N \parallel Sh_N)$  key-derivation followed by the ledger decryption/spend check — and *nothing* more: the search is over the 96-bit  $Sh_N$  *directly*, so Shamir reconstruction (the holder’s step of interpolating  $Sh_N$  from  $r_N$  recognized points) never enters the oracle loop. One quantum attempt is thus KDF + decryption, not Shamir + KDF + decryption.

no incremental sanction has no reason to stop short of the worse act, with *witness elimination* as its standard illustration. Read at that level of abstraction, the Deadly-Race Lemma is a witness-elimination argument in which the thing the victim retains is not testimony but a claim on a rivalrous prize. Independent arrival at the same mechanism from an unrelated field corroborates it; we claim no priority over either literature.

What is worth stating precisely is *where* our setting departs from Selten’s, because the departure is the reason our conclusion differs from his. Selten assumes the kidnapper’s threat to execute is credible *even though executing cannot improve his position* — and in classic kidnap-for-ransom that assumption is correct, because the prize is a ransom transferred by a third party and is not contestable once paid. The hostage’s survival threatens nothing the kidnapper holds. Bitcoin self-custody removes exactly that property. The prize is rivalrous and bearer, and a design that leaves  $\mathcal{V}$  a feasible path leaves it *contestable after release*: a surviving  $\mathcal{V}$  may spend first and take  $W$ . In Selten’s terms our  $(1 - p_A)$  is identically zero; the Deadly Race is his game with a non-degenerate post-release contest, and it is that single structural change that turns execution from a threat instrument into an action with positive expected value. The contribution here is therefore not the mechanism but its locus: marginal deterrence studies how a *legislator’s* penalty schedule moves this margin, whereas we show that a *custody designer* moves the same margin, and that a widely shipped class of designs moves it the wrong way.

Two qualifications follow. First, the counterweight is real and we do not minimize it: Iqbal, Masson and Abbott [50] extend Selten precisely by letting the capture probability rise when the hostage is executed, which is our  $c$  made explicit. Our claim is that the gain side can exceed  $c$  for large  $W$ , not that  $c$  is small — and Lemma 1 is stated as a threshold condition for exactly that reason. Second, the deterrence literature in this vein — Lapan and Sandler on precommitment to non-negotiation [51] — addresses whether the encounter *occurs*. Our results concern the game once it has, and we are careful to keep the two apart wherever a design claims deterrence as a structural property (§5).

## 8 Limitations and scope

---

**The tacit premise is unproven — the softest assumption.** Everything on the tacit route (hence “TGPO keeps all Four Properties”) rests on the existence of a *perceptual oracle*: a stimulus a human can *recognize* but not *articulate*. For TKBA it must simultaneously (i) admit a *binary encoding*; (ii) be *feasible to memorize* at the stipulated minimum ( $\geq 96$  bits); and (iii) be *infeasible to transmit within the coercion-time budget  $t$*  (Definition 12). None of the three is proved here; their conjunction is an empirical, psychometric claim awaiting a field study. Perturbed Burning-Ship fractal locations are our current instantiation — the best present candidate, with the design record giving the grounds for expecting them to hold — but until a study establishes such an oracle, **the validity of TKBA remains an assumption, not a result**, and it is the softest premise in the paper. Its imperfection is precisely the residual window (§4.6), where transmission becomes possible near the end of the derivation.

**Memorization cost, and why it is affordable.** A separate objection to a brain-only setup is not whether the oracle *exists* but whether living with it is too costly — carrying  $\geq 96$  bits of tacit recall indefinitely. Two things bound that cost. First, retention is *assisted*: a spaced-repetition companion tool [52] schedules review (an SM-2 / Anki-class algorithm), turning a one-off feat of memory into a light, decaying maintenance load. Second, the feat is historically ordinary: before mobile computation offloaded it, ordinary people routinely held a handful of phone numbers, several street addresses, and a working mental map of their city’s grid — recall comparable to what TKBA asks. The cost is real but bounded, and it is the price of a secret no device can be seized to reveal.

**Adoption precedent — demanding standards clear once the pain is legible.** The cost objection also has an empirical form — *users will never accept this much ceremony* —

and it met a natural experiment in the July 2026 seed-entropy incident (§9). Per the vendor’s own advisory, seeds created with 50-plus fair, independent dice rolls were not at risk from the defect [53]; within days, community guidance converged on exactly that manifestly labor-intensive procedure as the default counsel. A second and independent datum sits in the market for the rival method: holders pay for consulting that builds a decoy, and rehearse the story until it is believable (Remark 33). Both data are anecdotal, and neither burden is TKBA’s — performing a lie is not carrying 96 bits of tacit recall — but together they locate the asymmetry. TKBA’s efficacy is conjectural, resting on a perceptual-oracle premise no study has established; the decoy’s is not conjectural but condemned, in theory by Kerckhoffs and NNPP and in the record by the registry (§5.6) — a record obscurity itself thins, which is why the condemnation has to rest on the theory (Remark 37). Effort tolerance, in other words, is not a constant of the population; it reprices overnight once a threat is legible, and it is already being spent — on the one of the two that theory and the record jointly condemn. TKBA’s memorization load stands to the wrench-attack class as dice-rolling stands to entropy failure — demanding, verifiable by the holder alone, and adoptable at scale when the danger is understood — and making the danger legible *before* the corresponding incident is precisely what a formal bar is for.

**Stipulations aligned with standard practice.** Two quantities are stipulated, not derived. (a) The floor rests on the feasibility line S1, S2, S3 (Stipulation 1) — a standard *key-size* stipulation ( $2^{64}$  memory-hard and  $2^{96}$  cheap evaluations are beyond reach, in line with accepted security-level conventions); we do *not* claim a reduction to a named hard problem, and none is needed for a key-sizing argument. (b) The coercion-time budget  $t$  is a stipulation from *kidnapping-duration statistics* — read off the physical-attack registry [2], with the common-sense monotonicity that a kidnapping’s material feasibility *falls* with its duration. A *long-hold* adversary who sustains coercion past the time-gate ( $t > G$ ) is therefore out of scope: the tacit route assumes  $t < G$ , and against unbounded patience no in-head time-gate suffices. Most recorded attacks are in fact short hand-over-now events (§5.6).

**Single-shot utility, and why it is the conservative case.** The kill decision is modeled once, risk-neutrally. Risk-neutrality is *worst-for-victim*: a risk-averse  $\mathcal{A}$  discounts the uncertain kill-gain  $(1 - p_A)W$  and so kills *less* readily, while the cost side is already carried by  $c$  — so the single-shot risk-neutral model is a fortiori safe. *Repeated* encounters are absorbed by the construction’s re-key pathway: a victim who survives with  $\mathcal{K}$  uncompromised — itself a spectacular outcome, achieved with no delegate and no geographic split — has every incentive to harden physically for the short time needed to *re-key* the stash to a fresh setup (the level-up / shallow-stage-retirement pathway). *Inheritance* we leave beyond scope except for its design target: the protocol should render *assassinate-then-coerce-heir* non-viable — simplest remedy, a dead-man switch with heirs *also* TGPO-protected, so coercing an heir faces the same floor.

**The one-line takeaway.** Under TGPO — or a geographic setup, available only to those who can hold multiple independent sites — you *survive a WDY attack with individually custodied  $\mathcal{K}$  intact*. Alternatively you resort to shared custody, in which case, *unless* your stash is *demonstrably* more available to some remote (trusted) delegate than to you, the outcome of a wrench attack is “gray” — hence, by the No-Gray-Area Criterion, a sufficiently vicious attacker, or a high enough stash ( $W > c/\pi$ ), gets you killed.

**Composition across level-up.** DRL-safety across the setup pathway (re-keys, retirement-via-chain, a second wrench attacker who effectively starts a stage deeper) is a *deduction* from the per-encounter results applied at each stage, not a fresh theorem; we note it without re-proving it, and leave a full multi-encounter treatment to the utility-model extension above.

**Also out of scope.** Adversarial-agent scenarios (testator/heir mutual distrust, malicious interference with rotation); device side-channels (timing, EM, acoustic) during active tacit recall; long-horizon cryptographic erosion; and mass on-chain correlation beyond standard taproot / private-channel privacy. Finally, the framework fixes a *bar* (Criterion 1); it does not claim TGPO is the unique construction that clears it, and the residual window (§4.7) is a genuine bounded sliver argued incentive-safe rather than eliminated.

## 9 Ethics and responsible disclosure

---

This is offensive-flavored security research with direct life-safety stakes: we formally derive an *assassination* incentive (Lemma 1). We therefore state plainly why publishing it is, on balance, protective.

**The incentive is disclosed, not created.** The Deadly Race is a property of a situation a coercer already faces; naming it hands nothing to an attacker who, by Kerckhoffs (Assumption 1), is already assumed to know the protocol and the landscape. What publication changes is only that *defenders* can now measure a design against the bar (Criterion 1) and reject the ones that manufacture the incentive. We describe an incentive and the failure modes of existing defenses — not a targeting method, and no operational how-to; the paper’s constructive content is defensive.

**Obscurity is not a sound reason to withhold.** The alternative — not publishing analyses of attack vectors and the failure modes of deployed defenses — is itself a bet on obscurity, and both theory and data say that bet loses. As Kerckhoffs’ principle has held for a century, a defense that works only while the adversary stays ignorant provides something worse than defenselessness: *a false sense of security*. Empirically, the physical-attack registry [2] and the accompanying economic analysis [8] show exactly what the principle predicts — wrench attacks are frequent and rising, and the community’s entrenched counsel (*keep a low profile, never disclose ownership, deny if asked* — advice prominent channels actively promote, cataloged in [8]) has not stemmed them. Security knowledge diffuses to criminal actors regardless of academic reticence; withholding it handicaps only the defenders who would otherwise design against the very incentive we name. Publishing the Deadly Race, *together with* a bar that rejects the designs which create it, is net-protective of the holders it concerns.

**The quiet years compound — a supply-side precedent.** The ecosystem has already run this experiment on itself, on the *generation* side of the key lifecycle. A March-2021 firmware defect left a fallback software PRNG silently live as the production entropy path of a major hardware wallet (Coldcard) [54, 53]; the weakness was invisible to inspection — pseudorandom output passes every eyeball test, which is why entropy health cannot be observed from outside — and it stayed quiet for over five years, until a first wave of sweeps drained  $\approx \$70\text{M}$  from  $\approx 1,200$  addresses in 41 minutes on 30 July 2026 [55]. The episode teaches two things, and both transfer to coercion verbatim. First, the quiet years were not obscurity *working*: the loss was contracted the moment the defect shipped, and every silent month enrolled more holders into the eventual sweep — exposure compounding while it read, to everyone, as safety. Second, there was no “quiet fix” even in principle: a warning channel wide enough to reach the holders is wide enough to reach the attacker, so disclosure and attack start the same clock, and the only tenable posture is a design that survives full disclosure (the vendor’s prompt patch — fixed firmware for every affected model within two days — and candid post-mortem being the honest version of exactly that). Wrench-resistance sits in the same epistemic position as entropy health: a holder might not realise from outside whether their custody design protects them or endangers them. The community’s collective silence about deadly-race dangers — the norm against saying plainly that feasible-but-unfinished recovery paths endanger their owners — compounds exposure the same way the quiet years did, while it, too, reads as prudence. Naming the incentive is the disclosure; the quiet alternative is how it gets worse.

**Framing and restraint.** We keep the framing on protecting at-risk holders, disclose our own bounded residual honestly (§4.6) rather than overclaim, and treat the shipping alternatives we classify (§5) as legitimate clearances by the other route — not targets.

## 10 Conclusion

---

The Deadly Race turns “graceful degradation” from a virtue into a liability whenever the degraded state is a feasible-but-unfinished path to the secret: by Lemma 1, that state is a race, and the

race pays an assassination incentive. The No-Gray-Area Criterion is the resulting design bar, and the reachability corollary shows there are exactly two honest ways to clear it. The TGPO orbit clears it by removing the race entirely — a floor that leaves  $\mathcal{K}$  prohibitive to reach from every seizable state bar a bounded residual (Theorem 3, §4.6) — and so keeps custody individual, whereas time-lock-rescue vaults clear the same bar by removing the attacker’s reach at the cost of a trusted third party.

**What devicelessness buys: a threat model that stays simple.** Beyond clearing the bar, keeping the whole setup in the head keeps a series of otherwise-open cans of worms firmly shut — and it is this that lets the argument above stay *a fortiori*, granting  $\mathcal{A}$  every advantage and still finding no gray state, while keeping the game single-shot rather than a tangle of sub-games. The tacit, deviceless TGPO holder enjoys:

1. a **minimal window** in which  $\mathcal{K}$  is seizable at all — the bounded residual (§4.6), not a standing device (*the funds survive*);
2. a *separate* **minimal window** of deadly-race exposure (*the owner survives*) — distinct from (1), and only mostly overlapping it: once the second-to-last point of stage  $N$  is surrendered, the stash is effectively lost while the deadly-race window is at its *peak* (an impatient  $\mathcal{A}$ , or one perceiving a threat to its escape, may still kill);
3. **no hostage device** — and so none of the game-theoretic bluffing maze of ransoming a seized token (§3.7);
4. an NNPP that is a **dead end**:  $\mathcal{A}$  may assume an undisclosed shortcut device exists (Assumption 4), but killing  $\mathcal{V}$  cannot extract it, and  $\mathcal{A}$  knows  $\mathcal{V}$  has no *need* of one — so denial is inherently plausible (Remark 20);
5. **no trusted third party** (route R1, not R2);
6. **maximum portability** — no geographic binding to sites of custody (Definition 11);
7. the **lowest operational cost** of the three routes.

Each is a simplification that device-, geography-, or delegate-dependent designs forfeit; taken together they are why the deviceless secret admits so plain a threat model.

**Open items for the full draft.** The impossibility chain is drafted here (§3); the remaining work is: an *explicit adversary game* —  $\mathcal{A}$  a PPT adversary augmented by a bounded physical-coercion oracle — making Criterion 1 a theorem rather than a stipulated bar; a utility model richer than the single-shot  $(1 - p_A)W > c$  (repeated encounters, partial extraction, heir cascades,  $\mathcal{A}$ ’s risk-aversion); the floor *reduction* discharging §4’s  $p_A = 0$  obligation; a quantitative  $\mathbf{H}^*$  per-evaluation calibration tying S1/S2 to hardware; and an empirical grounding of the residual-window argument against the wrench-attack registry [2].

**Availability.** TGPO is implemented in an open-source reference application (<https://github.com/Yuri-SVB/Great-Wallet>); the protocol’s public design documentation and the registry-coding script behind §5.6 are maintained at <https://github.com/Yuri-SVB/great-wall-docs>.

## A A game-based formulation

---

The body states coercion-safety in a *decision-theoretic* register: the coercion game **Coerce<sub>II</sub>** (Definition 5) asks whether a *priced* adversary  $\text{WDY}[W, c]$  ever strictly prefers to kill. Readers from the cryptographic tradition will expect the complementary *experiment-and-advantage* form. We give it here — and, mirroring the two axes the residual splits into (§4.6 for funds, §4.7 for the deadly race), we give *two* experiments over one shared WDY interaction: a **seizure game**, whose winning event is the familiar “the adversary reaches  $\mathcal{K}$ ,” and a **deadly-race game**, whose winning event is the novel “the adversary reaches a *gray* state” — the one that manufactures the

kill incentive (Lemma 1), and which an adversary can trigger even when it provably cannot reach  $\mathcal{K}$ .

**Two adversaries, and why the primed one is the worst case.** The seizure game keeps the body’s adversary  $\mathcal{A}$  (holder  $\mathcal{V}$ ): it wants the *money*, seeks  $\mathcal{K}$ , and would kill only *instrumentally*, where the priced calculus  $(1 - p_A)W > c$  makes a kill raise its take (Definition 4). Its question is fund-safety — *does  $\mathcal{A}$  reach  $\mathcal{K}$ ?* The deadly-race game instead posits a *hypothetical* adversary  $\mathcal{A}'$  (holder  $\mathcal{V}'$ , primed to keep the two games apart) that wants the deadly race *for its own sake*: it plays to reach a gray state. This is deliberately the worst case for the kill question. Reaching a gray state is *necessary* for any kill to pay (Lemma 1); so if even the DR-hunting  $\mathcal{A}'$  cannot reach one — the coercion-resistance result below,  $\Pr[\text{win}] = 0$  — then the money-seeking  $\mathcal{A}$ , who would kill only instrumentally, never lands in a state where killing  $\mathcal{V}$  pays.  *$\mathcal{A}'$  trying and failing to manufacture a deadly race is exactly what certifies that  $\mathcal{A}$  will not run one.* Throughout,  $p_A$  and  $\pi$  stay unprimed: they are the body’s state functionals (Definitions 2, 3), read off the resulting position, and each game is a different *condition* on them.

### A.1 The shared interaction

**Definition 13** (The WDY interaction). Fix a custody protocol  $\Pi$ , a WDY adversary  $\mathcal{X}$  (Assumption 2), and the honest holder  $\mathcal{V}$ .

1. **Setup.** The challenger runs  $\Pi$ ’s honest setup for  $\mathcal{V}$ , producing the spend authority  $\mathcal{K}$  (§3.1), the co-located device and explicit-secret state  $D$  (vault and time-lock blobs, commitment scripts, review schedules), and the tacit state  $T$  (the holder’s recognition capability). Public parameters  $\text{pp}$ , including the salt  $\sigma$ , are given to  $\mathcal{X}$  (clause D0).
2. **Interaction.** Within a single attack site,  $\mathcal{X}$  queries, in any order:
  - **Seize** — returns the co-located device and explicit-secret state  $D$  (clauses D4,D5).
  - **Coerce( $t$ )** — for a coercion budget  $t$ , returns every value  $\mathcal{V}$  is *able* to transmit under full cooperation (clause D3); the tacit state  $T$  is available only as an *evaluate-only* oracle the challenger runs on  $\mathcal{V}$ ’s behalf, never as a transmissible value (clauses D6,D7).
  - **Compute** — any computation on the feasibility line (clause D1): at most  $2^{64}$  cheap and strictly fewer than  $2^{64}$  memory-hard evaluations.
 Any remote delegate is not coercible through  $\mathcal{V}$  (clause D8).
3. **Resulting position.** Let  $R$  be the resulting set of racers,  $p_A = p_A(R)$  the probability  $\mathcal{X}$  then reaches  $\mathcal{K}$  (Definition 5), and  $\pi = p_A(R \setminus \{\mathcal{V}\}) - p_A(R)$  the holder’s pivotality. Each game below is a predicate on this  $(p_A, \pi)$ .

### A.2 The seizure game — fund-safety

**Definition 14** (Seizure experiment; seizure-resistance). The **seizure experiment**  $\text{Sz-Exp}_\Pi(\mathcal{A})$  runs the interaction (Definition 13) with adversary  $\mathcal{X} = \mathcal{A}$  and holder  $\mathcal{V} = \mathcal{V}$ , and outputs 1 (“ $\mathcal{A}$  seizes”) iff  $\mathcal{A}$  feasibly reaches  $\mathcal{K}$  — equivalently  $p_A > 0$  (Definition 2).  $\Pi$  is **seizure-resistant** if

$$\Pr[\text{Sz-Exp}_\Pi(\mathcal{A}) = 1] = 0$$

for every WDY adversary  $\mathcal{A}$ .

This is the fund-safety / seizability axis — the game face of the *floor*.

**Corollary 6** (TGPO is seizure-resistant up to the residual). *The Time-Gated Perceptual Oracle (§4) is seizure-resistant except on the residual seizure window (§4.6).*

*Proof.* By Theorem 3, from every seizable state save that window  $K_N$  meets the floor, so no feasible path to  $\mathcal{K}$  exists and  $p_A = 0$ ; hence  $\text{Sz-Exp}$  outputs 0. On the residual seizure window  $\text{Sz-Exp}$  may output 1 — the  $o_N$ -live coincidence where  $\mathcal{A}$  does reach  $\mathcal{K}$ ; the loss there is of *funds*, bounded as argued in §4.6.  $\square$

### A.3 The deadly-race game — deadly-race-safety

**Definition 15** (Deadly-race experiment). The **deadly-race experiment**  $\text{DR-Exp}_\Pi(\mathcal{A}')$  runs the interaction (Definition 13) with adversary  $\mathcal{X} = \mathcal{A}'$  and holder  $\mathcal{Y} = \mathcal{V}'$ , and outputs 1 (“ $\mathcal{A}'$  wins”) iff  $\pi > 0$ ; for the individual-custody case  $R = \{\mathcal{V}'\}$  this is exactly a *gray* outcome,  $0 < p_A < 1$  (Definition 3).

**Definition 16** (Coercion resistance).  $\Pi$  is **coercion-resistant** if

$$\Pr[\text{DR-Exp}_\Pi(\mathcal{A}') = 1] = 0$$

for every WDY adversary  $\mathcal{A}'$ .

*Remark 39* (Where the two games part — the clean loss). The two experiments read the same position  $(p_A, \pi)$  but score it differently, and they diverge at exactly one point. At  $p_A = 0$  both output 0 (no feasible path: the defender clears both). On the *gray* interval  $0 < p_A < 1$  both output 1 ( $\mathcal{A}$  holds a path and  $\mathcal{V}'$  is pivotal: the defender fails both). They part only at *full compromise*  $p_A = 1$ : the seizure game outputs 1 (the stash is gone) while the deadly-race game outputs 0 ( $\pi = 0$  — with the adversary already certain to win, removing the holder changes nothing, so no kill incentive arises). That single disagreement is the whole content of “*a clean loss is deadly-race-safe*” (Remark 9): losing the funds outright is a seizure-game defeat but a deadly-race-game *non-event*. Conversely — the paper’s negative result restated — an adversary that provably *cannot* win the seizure game (cannot reach  $\mathcal{K}$ ; e.g. one facing a graceful-recovery vault it cannot open within the custody window) can still win the *deadly-race* game by reaching a gray state. The deadly-race axis is therefore *not* implied by fund-safety; it is a separate axis, exactly as §4.6 and §4.7 are separate. Fund-safety the construction handles on its own terms — *on top of* removing the gray race, it minimizes the window in which the stash is seizable at all, via the wipe-before-memory-hard step of (4) and the floor (Theorem 3).

*Remark 40* (The residual window: from 0 to  $\varepsilon$ ). Definition 16 is the information-theoretic, no-gray-area form. The residual window (§4.7) replaces the exact 0 by a bounded  $\varepsilon = \Pr[\text{Corr} \cap \text{Esc}]$ : call  $\Pi$   **$\varepsilon$ -coercion-resistant** if  $\Pr[\text{DR-Exp}_\Pi(\mathcal{A}') = 1] \leq \varepsilon$  for all  $\mathcal{A}'$ . The claims below are the  $\varepsilon = 0$  statements; the residual makes them  $\varepsilon$ -small — not zero — in the sliver argued incentive-safe there. In plain terms, this sliver is the moment the attack coincides with a *live* final orbit point  $o_N$  — currently or imminently materialized — and  $\mathcal{A}'$  prevents the wipe (or reset) that would otherwise close it; only then is transmissible material briefly within reach, which is why the window is a bounded  $\varepsilon$  rather than a standing exposure.

### A.4 Consistency with the main results

**Theorem 4** (The experiment coincides with the Criterion).  $\Pi$  is *coercion-resistant* (Definition 16) if and only if  $\pi = 0$  at every reachable state — equivalently, iff  $\Pi$  satisfies the *No-Gray-Area Criterion* (Criterion 1), iff  $\Pi$  is *coercion-safe* in the decision-theoretic game (Definition 5).

*Proof.* DR-Exp outputs 1 exactly on a reachable state with  $\pi > 0$ , so  $\Pr[\text{win}] = 0$  iff no reachable state is pivotal. By Theorem 1,  $\pi = 0$  everywhere is equivalent to coercion-safety and to Criterion 1. The two clearance routes of Corollary 2 — individual custody (R1) and a remote un-coercible racer (R2) — are exactly the two ways to force  $\Pr[\text{win}] = 0$ .  $\square$

**Corollary 7** (TGPO is coercion-resistant). The *Time-Gated Perceptual Oracle* (§4) is *coercion-resistant*.

*Proof.* The experiment outputs 1 only on a *gray* state ( $\pi > 0$ ); both clear endpoints  $p_A \in \{0, 1\}$  give  $\pi = 0$  and hence output 0. By Theorem 3 every state reachable by **Seize**, **Coerce**, and **Compute** leaves  $\mathcal{K}$  at the floor, so the reachable endpoint is the *clear failure*  $p_A = 0$  — the opposite endpoint  $p_A = 1$  would require feasibly reaching  $\mathcal{K}$ , which the same theorem bars. Either way  $\pi = 0$ , so the

experiment never outputs 1 and  $\Pr[\text{win}] = 0$ . The only place  $p_A$  rises above 0 (to 1 in the worst instant) is the residual window (§4.6), where the tacit premise momentarily fails; accounting for it yields  $\varepsilon$ -coercion resistance.  $\square$

*Remark 41* (Reading the classification as adversaries). Under Definition 16, the classification of §5 is a statement about which designs admit a winning  $\mathcal{A}'$ . Against a time-lock-rescue vault advertised as *individual* custody, the adversary that calls **Seize** and then waits out the delay is winning: the rescue racer is co-located and reachable (clause D4), so  $0 < p_A < 1$  and  $\pi > 0$ . Routing rescue to a genuinely remote, un-coercible delegate removes the reachable racer and restores  $\Pr[\text{win}] = 0$  — the honest R2 case, at the cost of individual custody.

## References

---

- [1] Danny Dolev and Andrew C. Yao. On the security of public key protocols. *IEEE Transactions on Information Theory*, 29(2):198–208, 1983.
- [2] Jameson Lopp. Known physical bitcoin attacks. <https://github.com/jlopp/physical-bitcoin-attacks>, 2025. Community registry of documented “\$5 wrench” incidents. Coded at commit 9a4a62a (2026-08-13), the state snapshotted in `registry-analysis/`; the registry is live and later revisions do not reproduce these figures.
- [3] Marilyne Ordekian, Gilberto Atondo-Siu, Alice Hutchings, and Marie Vasek. Investigating wrench attacks: Physical attacks targeting cryptocurrency users. In Rainer Böhme and Lucianna Kiffer, editors, *6th Conference on Advances in Financial Technologies (AFT 2024)*, volume 316 of *LIPICs*, pages 24:1–24:24. Schloss Dagstuhl — Leibniz-Zentrum für Informatik, 2024.
- [4] Auguste Kerckhoffs. La cryptographie militaire. *Journal des sciences militaires*, IX:5–38, 1883.
- [5] Jameson Lopp. Home invasion and torture of streamer RATIRL (sweden, 13 january 2026). Entry in *Known Physical Bitcoin Attacks*, <https://github.com/jlopp/physical-bitcoin-attacks>, 2026. Registry incident: a League of Legends streamer who held no cryptocurrency was tortured for hours; attackers took only his Pokémon cards; 2 suspects arrested.
- [6] CertiK. Intel3d: H1 2026 wrench attacks. CertiK Intel3D report, <https://www.certik.com/certik-report/intel3d/intel3d-wrench-h1-2026>, 2026. 52 verified incidents in H1 2026 (39 in H1 2025); \$124.1M losses against \$10.5M; home invasions 20 of 52, up from 1; France 33 of 52; targeting driven by leaked databases, tax records and exchange data. Sources include French Interior Ministry statements and court documents, independent of the physical-attack registry; archived at <http://archive.today/2026.09.03-170208/https://www.certik.com/certik-report/intel3d/intel3d-wrench-h1-2026>.
- [7] Journal du Coin. France: une employée du fisc vendait des informations confidentielles. <https://journalducoin.com/actualites/france-employee-fisc-vendait-informations-confidentielles/>, 2026. Reported 2026-01-07: an employee of the French tax administration placed under judicial investigation for illicit consultation of the internal “Mira” database and transmission of the results; investigators report queries targeting cryptocurrency investors. Allegations under investigation, not a conviction; archived at <http://archive.today/2026.01.09-094000/https://journalducoin.com/actualites/france-employee-fisc-vendait-informations-confidentielles/>.
- [8] Yuri da Silva Villas Boas. Justification and economics: The security cost of bitcoin self-custody. Technical report, <https://github.com/Yuri-SVB/great-wall-docs/blob/main/justification-and-economics/JUSTIFICATION.pdf>, 2026.

- [9] Yuri da Silva Villas Boas. Coercion-resistant self-custody — threat model and design record. Technical documentation, [https://github.com/Yuri-SVB/great-wall-docs/blob/main/great-wallet/THREAT\\_MODEL.md](https://github.com/Yuri-SVB/great-wall-docs/blob/main/great-wallet/THREAT_MODEL.md), 2026.
- [10] Joël Alwen, Jeremiah Blocki, and Krzysztof Pietrzak. Depth-robust graphs and their cumulative memory complexity. In *Advances in Cryptology — EUROCRYPT 2017*, volume 10212 of *Lecture Notes in Computer Science*, pages 3–32. Springer, 2017.
- [11] Alex Biryukov, Daniel Dinu, and Dmitry Khovratovich. Argon2: New generation of memory-hard functions for password hashing and other applications. In *IEEE European Symposium on Security and Privacy (EuroS&P)*, pages 292–302, 2016.
- [12] Yuri da Silva Villas Boas. The four properties and tacit knowledge-based authentication (tkba). Technical documentation, <https://github.com/Yuri-SVB/great-wall-docs/blob/main/great-wallet/ARCHITECTURE.md> and <https://github.com/Yuri-SVB/great-wall-docs/blob/main/justification-and-economics/JUSTIFICATION.pdf>, 2026.
- [13] Yuri da Silva Villas Boas. Burning ship encoder — design document. Technical documentation, <https://github.com/Yuri-SVB/great-wall-docs/blob/main/great-wall-core/DESIGN.md>, 2026. Public-facing protocol and instantiation-parameter reference.
- [14] Yuri da Silva Villas Boas. A memorization-free, precomputation-resistant salt from the bitcoin timechain. Companion paper (Namtso), <https://github.com/Yuri-SVB/namtso-the-sacred-salt>, 2026.
- [15] Adi Shamir. How to share a secret. *Communications of the ACM*, 22(11):612–613, 1979.
- [16] Ronald L. Rivest, Adi Shamir, and David A. Wagner. Time-lock puzzles and timed-release crypto. Technical Report MIT/LCS/TR-684, MIT Laboratory for Computer Science, 1996.
- [17] Yuri da Silva Villas Boas. Flying turtle — outsourced anonymous timelock. Technical documentation, <https://github.com/Yuri-SVB/great-wall-docs/blob/main/great-wallet/ARCHITECTURE.md> (Flying Turtle module), 2026.
- [18] Rewind Bitcoin. Rewind bitcoin — timelock vaults with rescue. <https://rewindbitcoin.com/>, 2025. archived at <http://archive.today/2026.09.03-170016/https://rewindbitcoin.com/>.
- [19] BitVault. Bitvault — time-delayed bitcoin self-custody. <https://www.bitvault.sv/>, 2026. Vendor site. Product claims verified against the live site 2026-08-24; the “publicly verifiable, making kidnapping attempts pointless” formulation is no longer carried there and is treated in the text as a former formulation the vendor has acknowledged. Re-checked and archived 1 September 2026: <http://archive.today/2026.09.01-224759/https://www.bitvault.sv/%23how-it-works>.
- [20] Francesco Madonna. B-SSL: Bitcoin secure signing layer. Concept whitepaper (Apache-2.0), self-attributed to F. Madonna of BitVault SA and published by a GitHub account identifying itself with the BitVault organization, [https://github.com/ilghan/bssl-whitepaper/blob/main/B-SSL\\_WP\\_Oct\\_11\\_2025.pdf](https://github.com/ilghan/bssl-whitepaper/blob/main/B-SSL_WP_Oct_11_2025.pdf), repository archived at <http://archive.today/2026.09.01-225748/https://github.com/ilghan/bssl-whitepaper/>. The cited PDF is 352,036 bytes, SHA-256 64a0c9c4b440ae8ff90a7ba8ecb62bacef3acf153753b9240aee17c8723f5320, 2025. Dated 11 October 2025; not linked from the vendor’s own site, and the attribution could not be authenticated further. Specifies a 2-of-3 whose paths carry timelocks: a configurable 2 h–15 d relative CSV primary and two multi-year absolute-CLTV fallbacks.

- [21] Blockstream. Enable a duress PIN. Blockstream Help Center, <https://help.blockstream.com/blockstream-jade/add-more-security-functionality/enable-a-duress-pin>, 2026. Vendor documentation using “duress PIN” and “wallet-erase PIN” for the same erase feature: “an alternate PIN that you can enable to automatically delete the encrypted wallet information from Jade if entered.” Accessed 2026-08-17; archived at <http://archive.today/2026.09.03-164731/https://help.blockstream.com/blockstream-jade/add-more-security-functionality/enable-a-duress-pin>.
- [22] Blockstream. How do I set up a wallet-erase PIN? Blockstream Help Center, <https://help.blockstream.com/blockstream-jade/faqs/how-do-i-set-up-a-wallet-erase-pin>, 2026. Vendor documentation for the Jade wallet-erase PIN: an alternate PIN that deletes the recovery phrase from the device, described as “especially useful in the event of a physical threat,” displaying “Internal Error” on entry, and carrying the instruction to “make sure you have a backup before proceeding.” Accessed 2026-08-17; archived at <http://archive.today/2026.09.03-164738/https://help.blockstream.com/blockstream-jade/faqs/how-do-i-set-up-a-wallet-erase-pin>.
- [23] Coinkite. Settings menu — trick PINs. COLDCARD Documentation, <https://coldcard.com/docs/settings/>, 2026. Vendor documentation for the trick-PIN actions: *Brick Self*; archived at <http://archive.today/2026.09.03-164807/https://coldcard.com/docs/settings/>, *Wipe Seed*, *Look Blank*, and a *Duress Wallet* described as “a personal safety feature” — “If you must reveal a PIN under duress, give the duress PIN instead of your Main PIN.” Accessed 2026-08-24.
- [24] Coinkite. PIN entry — COLDCARD firmware documentation. <https://github.com/Coldcard/firmware/blob/7fc656a/docs/pin-entry.md>, 2026. Vendor firmware documentation. The duress wallet is derived from the real wallet by a one-way process; enabling the brick PIN warrants “triple-check your backups are correct”; and on the case where the decoy is not believed, “you may be asked what the actual duress PIN is, while under duress. We suggest providing the ‘brickme’ PIN in that case.” Accessed 2026-08-24.
- [25] SatoshiLabs. Create wipe code to erase device. Trezor Knowledge Base, <https://trezor.io/learn/security-privacy/personal-security-standards/create-wipe-code-to-erase-device>, 2026. Vendor documentation calling the feature a “self-destruct” PIN; archived at <http://archive.today/2026.09.03-164822/https://trezor.io/learn/security-privacy/personal-security-standards/create-wipe-code-to-erase-device> code which “erases all private data, including the wallet backup (recovery seed).” The page makes no mention of duress, coercion or physical threat. Accessed 2026-08-24.
- [26] brianddk. Add plausible deniability to wipe-code interface. Issue #2055, <https://github.com/trezor/trezor-firmware/issues/2055>, 2022. Opened 2022-01-10: requests a user-defined message on wipe-code entry so an erasure can be presented as a hardware fault under duress, rather than the device reporting that the wipe code was entered. Closed as *not planned*; archived at <http://archive.today/2026.09.03-164059/https://github.com/trezor/trezor-firmware/issues/2055>.
- [27] Blockstream. How does a passphrase work on Jade? Blockstream Help Center, <https://help.blockstream.com/blockstream-jade/faqs/how-does-a-passphrase-work-on-jade>, 2026. Vendor documentation for BIP-39; archived at <http://archive.today/2026.09.03-164822/https://help.blockstream.com/blockstream-jade/faqs/how-does-a-passphrase-work-on-jade> passphrases on Jade: the passphrase “is added to your Jade’s recovery phrase or SeedQR and is used to create your wallet,” “any number of passphrases can be entered to change the wallet you will log into,” and “your Blockstream Jade will not store your passphrase.” Notably, the page makes no claim of deniability, decoy or duress use — unlike the wallet-erase PIN pages. Accessed 2026-08-17.

- [28] Sparrow Wallet. Frequently asked questions. Sparrow Wallet documentation, <https://www.sparrowwallet.com/docs/faq.html>, 2026. Vendor documentation. Under “How can I choose where my wallet is saved?”, the `-d` data-directory flag is documented with the claim that “[t]his feature allows you to store all Sparrow data on removable media making for more plausible deniability.” Retrieved 2026-09-06, sha256 0e0273dd...4a839e; archived in <sources/web/sparrow-faq.html> and at <http://archive.today/2026.09.06-151741/https://www.sparrowwallet.com/docs/faq.html>.
- [29] Sparrow Wallet. Best practices. Sparrow Wallet documentation, <https://sparrowwallet.com/docs/best-practices.html>, 2026. Vendor guidance, tiered by holding size. At the first hardware-wallet tier: “The hardware wallet should be stored securely along with a backup of the seed words (the latter ideally in a different location).” Retrieved 2026-09-06; sha256 eb4acdf0...1e667af; archived in <sources/web/sparrow-best-practices.html> and at <http://archive.today/2026.09.06-152816/https://sparrowwallet.com/docs/best-practices.html>.
- [30] Tus Llaves, Tus Bitcoin. Servicios. Self-custody consultancy service listing, <https://tusllavestusbtc.com/#servicios>, 2026. paid tiers from \$100 to \$600 USD; lists “Passphrase (tu bóveda señuelo) — Billetera señuelo + billetera real. Protección ante coacción/coerción”. Accessed 2026-09-01; archived at <http://archive.today/2026.09.03-170027/https://tusllavestusbtc.com/%23servicios>.
- [31] Bitcoin Security Guide. Passphrases. <https://bitcoinsecurity.org/guide/passphrase/>, 2026. instructs the holder to keep a small decoy balance on the bare seed as “what you’d hand over under physical coercion”, and that “for this to work, the decoy balance must be believable”. Accessed 2026-09-01; archived at <http://archive.today/2026.09.03-170047/https://bitcoinsecurity.org/guide/passphrase/>.
- [32] Ong Bitcoin Argentina. Lunes de autocustodia — programa #2. Live self-custody programme, [https://www.youtube.com/live/X\\_cCwJubekQ](https://www.youtube.com/live/X_cCwJubekQ), 2026. archived at [http://archive.today/2026.09.03-171910/https://www.youtube.com/live/X\\_cCwJubekQ](http://archive.today/2026.09.03-171910/https://www.youtube.com/live/X_cCwJubekQ).
- [33] Ludwig von Mises. *Human Action: A Treatise on Economics*. Yale University Press, 1949. The source of *praxeology* as used here: economic theorems are deduced from the fact of purposive action rather than induced from observation, and history illustrates a theorem without being able to verify or falsify it.
- [34] Ignaz Semmelweis. *Die Ätiologie, der Begriff und die Prophylaxis des Kindbettfiebers*. C. A. Hartleben, Pest, Vienna and Leipzig, 1861. English edition: *The Etiology, Concept, and Prophylaxis of Childbed Fever*, trans. K. Codell Carter, University of Wisconsin Press, 1983.
- [35] K. Codell Carter and Barbara R. Carter. *Childbed Fever: A Scientific Biography of Ignaz Semmelweis*. Greenwood Press, 1994. Source for the reception history: the grounds on which the chlorine-washing result was rejected, and the open letters that followed the 1861 monograph.
- [36] Casa. Casa — collaborative multisig bitcoin custody. <https://casa.io/>, 2026. 2-of-3 / 3-of-5 multisig; client keys on separate devices plus a Casa co-signing key; archived at <http://archive.today/2026.09.03-165338/https://casa.io/>.
- [37] CertiK. Skynet wrench attacks report. CertiK Skynet report, <https://www.certik.com/skynet-report/skynet-wrench-attacks-report>, 2026. 72 verified physical-coercion incidents globally in 2025 (+75% year on year); \$40.9M confirmed losses, stated to be under-reported; Europe over 40% of incidents, France the highest national count; archived at <http://archive.today/2026.09.03-171000/https://www.certik.com/skynet-report/skynet-wrench-attacks-report>.

- [38] DL News. Ledger co-founder david balland and wife kidnapped in france. DL News, <https://www.dlnews.com/articles/regulation/ledger-cofounder-david-balland-and-wife-kidnapped-in-france/>, 2025. archived at <http://archive.today/2026.09.03-171044/https://www.dlnews.com/articles/regulation/ledger-cofounder-david-balland-and-wife-kidnapped-in-france/>.
- [39] Lov K. Grover. A fast quantum mechanical algorithm for database search. In *28th Annual ACM Symposium on Theory of Computing (STOC)*, pages 212–219, 1996.
- [40] Markus Grassl, Brandon Langenberg, Martin Roetteler, and Rainer Steinwandt. Applying Grover’s algorithm to AES: Quantum resource estimates. In *Post-Quantum Cryptography — 7th International Workshop (PQCrypto 2016)*, volume 9606 of *Lecture Notes in Computer Science*, pages 29–43. Springer, 2016.
- [41] Ari Juels, Dario Catalano, and Markus Jakobsson. Coercion-resistant electronic elections. In *Proceedings of the 2005 ACM Workshop on Privacy in the Electronic Society (WPES)*, pages 61–70. ACM, 2005.
- [42] Hristo Bojinov, Daniel Sanchez, Paul Reber, Dan Boneh, and Patrick Lincoln. Neuroscience meets cryptography: Designing crypto primitives secure against rubber hose attacks. In *21st USENIX Security Symposium (USENIX Security)*, pages 129–141, 2012.
- [43] Fabian Monrose and Aviel D. Rubin. Keystroke dynamics as a biometric for authentication. *Future Generation Computer Systems*, 16(4):351–359, 2000.
- [44] David Kahn. *The Codebreakers: The Story of Secret Writing*. Scribner, revised edition, 1996.
- [45] Ran Canetti, Cynthia Dwork, Moni Naor, and Rafail Ostrovsky. Deniable encryption. In *Advances in Cryptology — CRYPTO 1997*, volume 1294 of *Lecture Notes in Computer Science*, pages 90–104. Springer, 1997.
- [46] Dan Boneh, Joseph Bonneau, Benedikt Bünz, and Ben Fisch. Verifiable delay functions. In *Advances in Cryptology — CRYPTO 2018*, volume 10991 of *Lecture Notes in Computer Science*, pages 757–788. Springer, 2018.
- [47] Reinhard Selten. A simple game model of kidnapping. In Rudolf Henn and Otto Moeschlin, editors, *Mathematical Economics and Game Theory: Essays in Honor of Oskar Morgenstern*, volume 141 of *Lecture Notes in Economics and Mathematical Systems*, pages 139–156. Springer, Berlin, Heidelberg, 1977. The canonical game-theoretic model of kidnapping: a two-person sequential game between kidnapper and the hostage’s family, in which the threat to execute is assumed credible *even though the kidnapper cannot improve his situation by carrying it out*.
- [48] George J. Stigler. The optimum enforcement of laws. *Journal of Political Economy*, 78(3):526–536, 1970. Origin of marginal deterrence in the economic theory of crime: penalties must be graduated, or an offender already exposed to the maximum has no reason to stop short of the worse act.
- [49] Steven Shavell. A note on marginal deterrence. *International Review of Law and Economics*, 12(3):345–355, 1992. Formal treatment of marginal deterrence; witness elimination is its standard illustration — where the marginal sanction for a further killing is near zero, nothing deters it.
- [50] Azhar Iqbal, Virginie Masson, and Derek Abbott. Kidnapping model: An extension of Selten’s game. *Royal Society Open Science*, 4(12):171484, 2017. Relaxes Selten’s assumption that the capture probability is independent of the hostage’s fate, modelling the higher enforcement effort that follows an execution.

- [51] Harvey E. Lapan and Todd Sandler. To bargain or not to bargain: That is the question. *American Economic Review*, 78(2):16–21, 1988. Papers and Proceedings. On precommitment to no-negotiation policies and deterrence expenditure — the question of whether the encounter occurs, not of the game once it has.
- [52] Yuri da Silva Villas Boas. A spaced-repetition companion for setup recall. Technical documentation, <https://github.com/Yuri-SVB/great-wall-docs/blob/main/great-wall-t/ARCHITECTURE.md> (memorization-assistant module), 2026.
- [53] Coinkite Inc. Entropy technical backgrounder. Coldcard vendor post-mortem and advisory, <https://blog.coinkite.com/entropy-technical-backgrounder/>, 2026. Vendor’s candid account of the defect; fixed firmware for all affected models by 2026-08-01; seeds created with 50+ fair, independent dice rolls deemed not at risk; archived at <http://archive.today/2026.09.03-171027/https://blog.coinkite.com/entropy-technical-backgrounder/>.
- [54] Block, Inc. Predictable RNG fallback and 32-bit reseed in Coldcard firmware. Block Engineering post-mortem, <https://engineering.block.xyz/blog/predictable-rng-fallback-and-32-bit-reseed-in-coldcard-firmware>, 2026. Discovery write-up of the July 2026 seed-entropy incident: a preprocessor-guard defect, shipped in firmware v4.0.0 (2021-03-17), left a fallback software PRNG as the silent production entropy path; archived at <http://archive.today/2026.08.06-112407/https://engineering.block.xyz/blog/predictable-rng-fallback-and-32-bit-reseed-in-coldcard-firmware>.
- [55] CoinDesk. How bitcoin cold wallets lost \$70 million in an attack that never touched the devices. CoinDesk, <https://www.coindesk.com/tech/2026/08/01/how-bitcoin-cold-wallets-lost-usd70-million-in-an-attack-that-never-touched-the-devices>, 2026. First-wave figures: 1,082.65 BTC (about \$70M) drained from 1,196 addresses in a 41-minute window on 2026-07-30, per Galaxy Research’s mapping; archived at <http://archive.today/2026.08.01-211808/https://www.coindesk.com/tech/2026/08/01/how-bitcoin-cold-wallets-lost-usd70-million-in-an-attack-that-never-touched-the-devices>.